

▼ Name - Omkar Phansopkar

UID - 2022701007

D3 - CSE-DS

Dataset - Telco customer churn: IBM dataset

<https://www.kaggle.com/datasets/yeanzc/telco-customer-churn-ibm-dataset>

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)

```
!pip install branca plotly scikit-learn scipy imblearn --quiet
!pip install h3 folium --quiet
```

```
import numpy as np
from scipy import stats
from IPython.display import Image
import branca.colormap as cm
import pandas as pd
import seaborn as sns
import plotly.express as px
from plotly.offline import init_notebook_mode, iplot
import seaborn as sns
import matplotlib.pyplot as plt
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline
import h3
import matplotlib
import imblearn
import os
import folium
from sklearn.impute import SimpleImputer
```

▼ Data Description

CustomerID: A unique ID that identifies each customer.

Count: A value used in reporting/dashboarding to sum up the number of customers in a filtered set.

Country: The country of the customer's primary residence.

State: The state of the customer's primary residence.

City: The city of the customer's primary residence.

Zip Code: The zip code of the customer's primary residence.

Lat Long: The combined latitude and longitude of the customer's primary residence.

Latitude: The latitude of the customer's primary residence.

Longitude: The longitude of the customer's primary residence.

Gender: The customer's gender: Male, Female

Senior Citizen: Indicates if the customer is 65 or older: Yes, No

Partner: Indicate if the customer has a partner: Yes, No

Dependents: Indicates if the customer lives with any dependents: Yes, No. Dependents could be children, parents, grandparents, etc.

Tenure Months: Indicates the total amount of months that the customer has been with the company by the end of the quarter specified above.

Phone Service: Indicates if the customer subscribes to home phone service with the company: Yes, No

Multiple Lines: Indicates if the customer subscribes to multiple telephone lines with the company: Yes, No

Internet Service: Indicates if the customer subscribes to Internet service with the company: No, DSL, Fiber Optic, Cable.

Online Security: Indicates if the customer subscribes to an additional online security service provided by the company: Yes, No

Online Backup: Indicates if the customer subscribes to an additional online backup service provided by the company: Yes, No

Device Protection: Indicates if the customer subscribes to an additional device protection plan for their Internet equipment provided by the company: Yes, No

Tech Support: Indicates if the customer subscribes to an additional technical support plan from the company with reduced wait times: Yes, No

Streaming TV: Indicates if the customer uses their Internet service to stream television programming from a third party provider: Yes, No. The company does not charge an additional fee for this service.

Streaming Movies: Indicates if the customer uses their Internet service to stream movies from a third party provider: Yes, No. The company does not charge an additional fee for this service.

Contract: Indicates the customer's current contract type: Month-to-Month, One Year, Two Year.

Paperless Billing: Indicates if the customer has chosen paperless billing: Yes, No

Payment Method: Indicates how the customer pays their bill: Bank Withdrawal, Credit Card, Mailed Check

Monthly Charge: Indicates the customer's current total monthly charge for all their services from the company.

Total Charges: Indicates the customer's total charges, calculated to the end of the quarter specified above.

Churn Label: Yes = the customer left the company this quarter. No = the customer remained with the company. Directly related to Churn Value.

Churn Value: 1 = the customer left the company this quarter. 0 = the customer remained with the company. Directly related to Churn Label.

Churn Score: A value from 0-100 that is calculated using the predictive tool IBM SPSS Modeler. The model incorporates multiple factors known to cause churn. The higher the score, the more likely the customer will churn.

CLTV: Customer Lifetime Value. A predicted CLTV is calculated using corporate formulas and existing data. The higher the value, the more valuable the customer. High value customers should be monitored for churn.

Churn Reason: A customer's specific reason for leaving the company. Directly related to Churn Category.

Source This dataset is detailed in: <https://community.ibm.com/community/user/businessanalytics/blogs/steven-macko/2019/07/11/telco-customer-churn-1113>

Downloaded from: <https://community.ibm.com/accelerators/?context=analytics&query=telco%20churn&type=Data&product=Cognos%20Analytics>

There are several related datasets as documented in: <https://community.ibm.com/community/user/businessanalytics/blogs/steven-macko/2018/09/12/base-samples-for-ibm-cognos-analytics>

```
data = pd.read_excel('/content/drive/MyDrive/datasets/fds_customer_churn/Telco_customer_churn.xlsx')
data.sample()
```

	CustomerID	Count	Country	State	City	Zip Code	Lat	Long	Latitude
1565	4192-GORJT	1	United States	California	Fremont	94555	37.555473,	-122.080312	37.555473

1 rows x 33 columns

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 33 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            7043 non-null   object
1   Count                 7043 non-null   int64
2   Country               7043 non-null   object
3   State                7043 non-null   object
4   City                 7043 non-null   object
5   Zip Code             7043 non-null   int64
6   Lat Long             7043 non-null   object
7   Latitude             7043 non-null   float64
8   Longitude            7043 non-null   float64
9   Gender               7043 non-null   object
10  Senior Citizen       7043 non-null   object
11  Partner              7043 non-null   object
12  Dependents          7043 non-null   object
13  Tenure Months        6995 non-null   float64
14  Phone Service        7043 non-null   object
15  Multiple Lines       7043 non-null   object
16  Internet Service     7043 non-null   object
17  Online Security      7043 non-null   object
18  Online Backup        7043 non-null   object
19  Device Protection    7043 non-null   object
20  Tech Support         7043 non-null   object
21  Streaming TV         7043 non-null   object
22  Streaming Movies     7043 non-null   object
23  Contract             7043 non-null   object
24  Paperless Billing    7043 non-null   object
25  Payment Method       7043 non-null   object
```

```

26 Monthly Charges    7043 non-null float64
27 Total Charges      7043 non-null object
28 Churn Label        7043 non-null object
29 Churn Value         7043 non-null int64
30 Churn Score         7043 non-null int64
31 CLTV                7043 non-null int64
32 Churn Reason        1869 non-null object
dtypes: float64(4), int64(5), object(24)
memory usage: 1.8+ MB

```

- what services customers use,
- type of contract
- the lifetime of the client in the service
- payment method
- the amount of monthly payments of customers and their total costs in the service,
- customer locations,
- gender and age of the client
- reason for churn (for clients in the churn)

```
data['Total Charges']
```

```

0      108.15
1      151.65
2       820.5
3     3046.05
4      5036.3
...
7038    1419.4
7039    1990.5
7040    7362.9
7041     346.45
7042    6844.5
Name: Total Charges, Length: 7043, dtype: object

```

```
data['Total Charges'] = pd.to_numeric(data['Total Charges'], errors='coerce')
data['Total Charges']
```

```

0      108.15
1      151.65
2       820.50
3     3046.05
4      5036.30
...
7038    1419.40
7039    1990.50
7040    7362.90
7041     346.45
7042    6844.50
Name: Total Charges, Length: 7043, dtype: float64

```

```
data.isnull().sum()
```

```

CustomerID      0
Count           0
Country         0
State           0
City            0
Zip Code        0
Lat Long        0
Latitude         0
Longitude        0
Gender           0
Senior Citizen  0
Partner         0
Dependents      0
Tenure Months   48
Phone Service   0
Multiple Lines  0
Internet Service 0
Online Security 0
Online Backup    0
Device Protection 0
Tech Support     0
Streaming TV     0
Streaming Movies 0
Contract         0
Paperless Billing 0
Payment Method   0
Monthly Charges  0
Total Charges    11
Churn Label      0

```

```
Churn Value      0
Churn Score      0
CLTV             0
Churn Reason     5174
dtype: int64
```

```
data.isnull().sum() / len(data)
```

```
CustomerID      0.000000
Count           0.000000
Country          0.000000
State           0.000000
City            0.000000
Zip Code        0.000000
Lat Long        0.000000
Latitude         0.000000
Longitude       0.000000
Gender           0.000000
Senior Citizen  0.000000
Partner         0.000000
Dependents      0.000000
Tenure Months   0.006815
Phone Service   0.000000
Multiple Lines  0.000000
Internet Service 0.000000
Online Security 0.000000
Online Backup   0.000000
Device Protection 0.000000
Tech Support    0.000000
Streaming TV    0.000000
Streaming Movies 0.000000
Contract        0.000000
Paperless Billing 0.000000
Payment Method  0.000000
Monthly Charges 0.000000
Total Charges   0.001562
Churn Label     0.000000
Churn Value     0.000000
Churn Score     0.000000
CLTV            0.000000
Churn Reason    0.734630
dtype: float64
```

```
data[data['Tenure Months'].isnull()]
```

	CustomerID	Count	Country	State	City	Zip Code	Lat	Long	Latitude	Longitude	Gender	...	Contract	Paper Bill
217	9944-HKVVB	1	United States	California	King City	93930	36.220761,	-120.980777	36.220761	-120.980777	Female	...	Month-to-month	
258	9524-EGPJC	1	United States	California	Walnut Creek	94595	37.862128,	-122.075197	37.862128	-122.075197	Female	...	Month-to-month	
332	6818-WOBHJ	1	United States	California	Fulton	95439	38.493888,	-122.777141	38.493888	-122.777141	Female	...	Month-to-month	
416	9391-EOYLI	1	United States	California	Lewiston	96052	40.704293,	-122.803899	40.704293	-122.803899	Male	...	Month-to-month	
419	3068-OMWZA	1	United States	California	Round Mountain	96084	40.923558,	-122.059933	40.923558	-122.059933	Male	...	Month-to-month	
438	2058-DCJBE	1	United States	California	Los Angeles	90012	34.065875,	-118.238728	34.065875	-118.238728	Male	...	Month-to-month	
515	9497-QCMMS	1	United States	California	Escondido	92026	33.21846,	-117.116916	33.218460	-117.116916	Male	...	Month-to-month	
568	3023-GFLBR	1	United States	California	San Bernardino	92408	34.084909,	-117.258107	34.084909	-117.258107	Female	...	Month-to-month	
573	9061-TIHDA	1	United States	California	Murrieta	92563	33.581045,	-117.14719	33.581045	-117.147190	Male	...	Month-to-month	
637	9221-OTIVJ	1	United States	California	Dunlap	93621	36.789213,	-119.140338	36.789213	-119.140338	Female	...	Month-to-month	
666	0407-BDJKB	1	United States	California	San Francisco	94114	37.758085,	-122.434801	37.758085	-122.434801	Male	...	Month-to-month	
669	3269-ATYWD	1	United States	California	San Francisco	94123	37.800254,	-122.436975	37.800254	-122.436975	Male	...	Month-to-month	
749	9992-RRAMN	1	United States	California	Riverbank	95367	37.734971,	-120.954271	37.734971	-120.954271	Male	...	Month-to-month	
824	5649-ANRML	1	United States	California	Clipper Mills	95930	39.562239,	-121.14836	39.562239	-121.148360	Male	...	Month-to-month	
830	8393-DLHGA	1	United States	California	Palermo	95968	39.435756,	-121.552071	39.435756	-121.552071	Male	...	Month-to-month	
842	5519-NPHVG	1	United States	California	Klamath River	96050	41.816595,	-122.948287	41.816595	-122.948287	Female	...	Month-to-month	
850	5498-TXHLF	1	United States	California	Whitmore	96096	40.637105,	-121.906949	40.637105	-121.906949	Female	...	Month-to-month	
937	0133-BMFZO	1	United States	California	Alhambra	91803	34.074736,	-118.145959	34.074736	-118.145959	Female	...	Month-to-month	
942	8623-TMRBY	1	United States	California	Chula Vista	91915	32.605012,	-116.97595	32.605012	-116.975950	Male	...	Month-to-month	
944	3208-YPIOE	1	United States	California	La Mesa	91942	32.782501,	-117.01611	32.782501	-117.016110	Male	...	Month-to-month	
945	2612-RANWT	1	United States	California	Mount Laguna	91948	32.830852,	-116.444601	32.830852	-116.444601	Female	...	Month-to-month	
956	3194-ORPIK	1	United States	California	San Diego	92113	32.697098,	-117.116587	32.697098	-117.116587	Female	...	Month-to-month	
1055	0618-XWMSS	1	United States	California	Wasco	93280	35.652242,	-119.4464	35.652242	-119.446400	Male	...	Month-to-month	
1056	5276-KQWHG	1	United States	California	Woody	93287	35.710244,	-118.881679	35.710244	-118.881679	Female	...	Month-to-month	
1066	2357-	1	United States	California	Altamont	91007	37.454924,	-122.451001	37.454924	-122.451001	Female	...	Month-to-month	

1099	COQEK	1	United States	California	Atherton	94027	37.454924, -122.203168	37.454924	-122.203168	Female	...	Month-to-month
1157	5449-FIBXJ	1	United States	California	San Jose	95127	37.375156, -121.795867	37.375156	-121.795867	Male	...	Month-to-month
1158	9837-BMCLM	1	United States	California	Stockton	95202	37.959706, -121.287669	37.959706	-121.287669	Male	...	Month-to-month
1244	5729-KLZAR	1	United States	California	Sacramento	95817	38.550722, -121.457314	38.550722	-121.457314	Female	...	Month-to-month
1246	8740-CRYFY	1	United States	California	Sacramento	95823	38.475465, -121.443625	38.475465	-121.443625	Male	...	Month-to-month
1475	6284-AHOOQ	1	United States	California	Avenal	93204	35.916943, -120.129921	35.916943	-120.129921	Male	...	Month-to-month
1480	4818-DRBQT	1	United States	California	Farmersville	93223	36.29878, -119.201028	36.298780	-119.201028	Male	...	Month-to-month
1484	4391-RESHN	1	United States	California	Lemon Cove	93244	36.462671, -118.997291	36.462671	-118.997291	Male	...	Month-to-month
1528	3500-RMZLT	1	United States	California	Fresno	93702	36.739385, -119.753649	36.739385	-119.753649	Female	...	Month-to-month
1534	3422-LYEPQ	1	United States	California	Carmel By The Sea	93921	36.554618, -121.922239	36.554618	-121.922239	Male	...	Month-to-month
1538	9507-EXLTT	1	United States	California	Seaside	93955	36.625114, -121.823565	36.625114	-121.823565	Female	...	Month-to-month

```

imputer = SimpleImputer(strategy='mean')
data['Tenure Months'] = imputer.fit_transform(data[['Tenure Months']])
data['Tenure Months'].isnull().sum()

```

0

```

data.groupby('Churn Label')['CustomerID'].nunique()

Churn Label
No      5174
Yes     1869
Name: CustomerID, dtype: int64

```

```
data[data['Total Charges'].isna()]
```

```

CustomerID Count Country State City Zip Code Lat Long Latitude
2234 4472-LVYGI 1 United California San 92408 34.084909, 34.084909
data.dropna(subset=['Country', 'State', 'City'], inplace=True)
2428 3115-C7M7D 1 United California Independence 92526 36.869584, 36.869584
data['calc_charges'] = data['Monthly Charges'] * data['Tenure Months']
2568 5700-1V0EQ 1 United California San Mateo 94401 37.590421, 37.590421
data.columns

Index(['CustomerID', 'Count', 'Country', 'State', 'City', 'Zip Code',
      'Lat Long', 'Latitude', 'Longitude', 'Gender', 'Senior Citizen',
      'Partner', 'Dependents', 'Tenure Months', 'Phone Service',
      'Multiple Lines', 'Internet Service', 'Online Security',
      'Online Backup', 'Device Protection', 'Tech Support', 'Streaming TV',
      'Streaming Movies', 'Contract', 'Paperless Billing', 'Payment Method',
      'Monthly Charges', 'Total Charges', 'Churn Label', 'Churn Value',
      'Churn Score', 'CLTV', 'Churn Reason', 'calc_charges'],
      dtype='object')

```

calculating difference between Total Charges and calculated charges

```

United 34.144702
data['diff_in_charges'] = data['Total Charges'] - data['calc_charges']
data['diff_in_charges']

0      0.45
1     10.25
2     23.30
3    111.65
4    -45.00
...
7038 -103.40
7039 -44.70
7040 -67.50
7041  20.85
7042 -128.40
Name: diff_in_charges, Length: 7043, dtype: float64

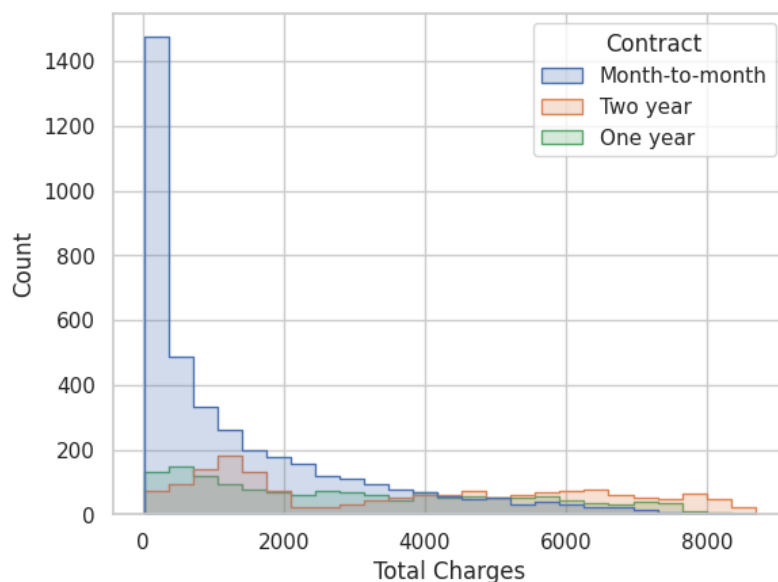
```

```

# Charges distribution by contract
sns.set(style="whitegrid")
sns.histplot(data=data, x="Total Charges", hue="Contract", element="step", common_norm=False)

```


<Axes: xlabel='Total Charges', ylabel='Count'>



```

# How are total charges distributed ?
plt.figure(figsize=(10, 6))
ax = sns.boxplot(data=data, x="Total Charges", color="red")
ax.set_xticks(range(0, 9000, 1000))
plt.subplots_adjust(left=0.1, right=0.9, top=0.9, bottom=0.1)

```



```
# Accounting irregularities?
data.groupby('Contract')[['Total Charges', 'diff_in_charges']].quantile([.50, .80, .90, .95])
```


		Total Charges	diff_in_charges	
Contract				
Month-to-month	0.50	679.5500	0.0000	
	0.80	2485.7300	24.6800	
	0.90	3844.0600	54.2000	
	0.95	4966.9200	87.3650	
One year	0.50	2657.5500	0.7750	
	0.80	5286.4600	55.0500	
	0.90	6341.2500	92.2000	
	0.95	7072.4725	133.3375	
Two year	0.50	3623.9500	0.5000	
	0.80	6399.2400	61.5300	
	0.90	7457.6100	97.5700	
	0.95	7922.3400	139.1800	

```
data['Total Charges'] = np.where(data['Total Charges'].isna() == True, data['calc_charges'], data['Total Charges'])
```

```
data.select_dtypes(include=['number']).columns
```

```
Index(['Count', 'Zip Code', 'Latitude', 'Longitude', 'Tenure Months',
      'Monthly Charges', 'Total Charges', 'Churn Value', 'Churn Score',
      'CLTV', 'calc_charges', 'diff_in_charges'],
      dtype='object')
```

```
data[['Total Charges', 'Churn Score',
      'calc_charges']].boxplot()
plt.title('Box Plots for Numerical Columns')
plt.ylabel('Value')
plt.show()
```





```
data = data.drop(['calc_charges', 'diff_in_charges'], axis=1)
```

```
# Glance of churn non churn
```

```
plt.pie(data['Churn Label'].value_counts(), labels=['No Churn', 'Churn'], autopct='%1.1f%%')
```

```
([<matplotlib.patches.Wedge at 0x7b729f523e50>,
 <matplotlib.patches.Wedge at 0x7b729f523d60>],
 [Text(-0.7393678277834757, 0.8144539368428056, 'No Churn'),
 Text(0.7393677515287918, -0.8144540060674139, 'Churn')],
 [Text(-0.4032915424273503, 0.44424760191425755, '73.5%'),
 Text(0.4032915008338864, -0.4442476396731348, '26.5%')])
```



```
data.groupby(['Country', 'State'])['CustomerID'].count()
```

```
Country      State
United States  California    7043
Name: CustomerID, dtype: int64
```

```
data['City'].nunique()
```

```
1129
```


```
# Geolocation scatter plot
```

```
fig = px.scatter_mapbox(data.groupby(['Latitude', 'Longitude'])['CustomerID'].count().reset_index(), lat="Latitude", lon="Longitude", mapbox_style="open-street-map")
fig.update_layout(mapbox_style="open-street-map")
fig.update_layout(margin={"r":0,"t":0,"l":0,"b":0})
fig.show()
```



```
# How many customers in each city ?
f, ax = plt.subplots(figsize=(18,5))
plt.bar(
    data.groupby('City')['CustomerID'].count().sort_values(ascending=False).head(
        10).index,
    data.groupby('City')['CustomerID'].count().sort_values(ascending=False).head(10).values,
)
ax.legend(fontsize = 14)
```

WARNING:matplotlib.legend.No artists with labels found to put in legend. Not
<matplotlib.legend.Legend at 0x7b729c7ddf00>



```
# Geolocation & number of clients
hex_level = 5

data['hex_id'] = data.apply(lambda x: h3.geo_to_h3(x['Latitude'], x['Longitude'], hex_level), axis=1)

hex_counts = data.groupby('hex_id')['CustomerID'].count().reset_index(name='total_clients')
hex_counts['center'] = hex_counts['hex_id'].apply(lambda x: h3.h3_to_geo(x))

color_range = [hex_counts['total_clients'].min(), hex_counts['total_clients'].max()]
colormap = cm.LinearColormap(["purple","red","orange","yellow","green"],vmin = min(color_range), vmax = max(color_range))


mean_lat, mean_lon = hex_counts['center'].apply(lambda x: x[0]).mean(), hex_counts['center'].apply(lambda x: x[1]).mean()
map_center = [mean_lat, mean_lon]
m = folium.Map(location=map_center, zoom_start=6, tiles='Stamen Terrain')

for _, row in hex_counts.iterrows():
    folium.Polygon(
        locations=h3.h3_to_geo_boundary(row['hex_id']),
        fill=True,
        fill_color=colormap(row['total_clients']),
        fill_opacity=0.7,
        stroke=False,
        tooltip=f"Number of clients: {row['total_clients']}"
    ).add_to(m)

colormap.caption = 'Number of clients'
m.add_child(colormap)

m
```

Make this Notebook Trusted to load map: File -> Trust Notebook



```
churn = data.assign(churn_clients = np.where(data['Churn Label']=='Yes',data['CustomerID'],None)).groupby(['hex_id']).agg({
```

```
data.groupby(['hex_id'])['CustomerID'].count()
```

```
hex_id
85280043ffffff      4
8528004bffffff      4
8528004ffffff        4
85280207ffffff      4
8528020bffffff      4
..
85485b03ffffff      5
85485b33ffffff      5
85485b63ffffff      5
85485babffffff      5
85485bb7ffffff     10
Name: CustomerID, Length: 714, dtype: int64
```

```
clients = data.groupby(['hex_id'])['CustomerID'].count().reset_index()
clients.sample(5)
```

	hex_id	CustomerID	
481	8529a227fffff	4	
162	8528305bfffff	8	
629	8529ab8bfffff	4	
573	8529a8b3fffff	4	
146	85283007fffff	8	

```
# Hex id for plotting purpose
churn_data = clients.join(churn.set_index(['hex_id']), on=['hex_id'])
churn_data
```

	hex_id	CustomerID	churn_clients	
0	85280043fffff	4	1	
1	8528004bfffff	4	1	
2	8528004fffff	4	1	
3	85280207fffff	4	1	
4	8528020bfffff	4	1	
...	
709	85485b03fffff	5	1	
710	85485b33fffff	5	1	
711	85485b63fffff	5	0	
712	85485babfffff	5	3	
713	85485bb7fffff	10	3	

714 rows x 3 columns

```
churn_data['churn_rate'] = churn_data['churn_clients']/churn_data['CustomerID']
churn_data['churn_rate']
```

```
0    0.25
1    0.25
```

```

2      0.25
3      0.25
4      0.25
...
709    0.20
710    0.20
711    0.00
712    0.60
713    0.30
Name: churn_rate, Length: 714, dtype: float64

```

churn_data

	hex_id	CustomerID	churn_clients	churn_rate
0	85280043ffffff	4	1	0.25
1	8528004bffffff	4	1	0.25
2	8528004ffffff	4	1	0.25
3	85280207ffffff	4	1	0.25
4	8528020bffffff	4	1	0.25
...
709	85485b03ffffff	5	1	0.20
710	85485b33ffffff	5	1	0.20
711	85485b63ffffff	5	0	0.00
712	85485babffffff	5	3	0.60
713	85485bb7ffffff	10	3	0.30

714 rows x 4 columns

```

churn_data['center'] = churn_data['hex_id'].apply(lambda x: h3.h3_to_geo(x))

color_range = [churn_data['churn_rate'].min(), churn_data['churn_rate'].max()]
colormap = cm.LinearColormap(["green","orange","red"],vmin = min(color_range), vmax = max(color_range))

mean_lat, mean_lon = churn_data['center'].apply(lambda x: x[0]).mean(), churn_data['center'].apply(lambda x: x[1]).mean()
map_center = [mean_lat, mean_lon]
m = folium.Map(location=map_center, zoom_start=6, width='100%', height='80%', tiles='Stamen Terrain')

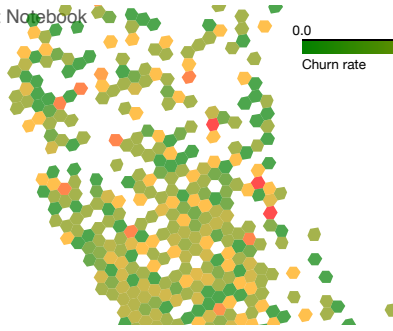
for _, row in churn_data.iterrows():
    folium.Polygon(
        locations=h3.h3_to_geo_boundary(row['hex_id']),
        fill=True,
        fill_color=colormap(row['churn_rate']),
        fill_opacity=0.7,
        stroke=False,
        tooltip=f"Churn rate: {row['churn_rate']}<br>Number of customers: {row['CustomerID']}"
    ).add_to(m)

colormap.caption = 'Churn rate'
m.add_child(colormap)

m


```

Make this Notebook Trusted to load map: File -> Trust Notebook




```
plt.scatter(
    churn_data['CustomerID'],
    churn_data['churn_rate'],
)
```

<matplotlib.collections.PathCollection at 0x7b729a390fa0>




```
fig, ax = plt.subplots()
ax.scatter(churn_data['CustomerID'], churn_data['churn_rate'], s=churn_data['churn_rate']*100, c=churn_data['churn_rate'],
ax.set_xlabel('Customer ID')
ax.set_ylabel('Churn Rate')
```

Text(0, 0.5, 'Churn Rate')



```
plt.hist(data['Tenure Months'], bins=50)
plt.legend()
plt.xlabel('Tenure Months')
plt.ylabel('Number of customers')
```

WARNING:matplotlib.legend.No artists with labels found to put in legend. Note that artists whose label start with an Text(0, 0.5, 'Number of customers')



```
data.groupby('Churn Label')['Tenure Months'].quantile([.50,.75,.90,.95])
```

```
Churn Label
No      0.50    38.0
        0.75    61.0
        0.90    71.0
        0.95    72.0
Yes     0.50    10.0
        0.75    31.0
        0.90    50.0
        0.95    60.0
Name: Tenure Months, dtype: float64
```

```
data.groupby('Churn Label')['Tenure Months'].mean()
```

```
Churn Label
No      37.578893
Yes     18.274855
Name: Tenure Months, dtype: float64
```

```
grouped = data.groupby(['Churn Reason'])['CustomerID'].count().reset_index().sort_values('CustomerID',
                                                                                          ascending=False)
```


```
grouped
```

```

Churn Reason CustomerID
1 Attitude of support person 192
# What is frequency of each reason for churning out ?
plt.bar(grouped['Churn Reason'], grouped['CustomerID'])
plt.xticks(rotation=90)

([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19],
[Text(0, 0, 'Attitude of support person'),
Text(1, 0, 'Competitor offered higher download speeds'),
Text(2, 0, 'Competitor offered more data'),
Text(3, 0, "Don't know"),
Text(4, 0, 'Competitor made better offer'),
Text(5, 0, 'Attitude of service provider'),
Text(6, 0, 'Competitor had better devices'),
Text(7, 0, 'Network reliability'),
Text(8, 0, 'Product dissatisfaction'),
Text(9, 0, 'Price too high'),
Text(10, 0, 'Service dissatisfaction'),
Text(11, 0, 'Lack of self-service on Website'),
Text(12, 0, 'Extra data charges'),
Text(13, 0, 'Moved'),
Text(14, 0, 'Limited range of services'),
Text(15, 0, 'Long distance charges'),
Text(16, 0, 'Lack of affordable download/upload speed'),
Text(17, 0, 'Poor expertise of phone support'),
Text(18, 0, 'Poor expertise of online support'),
Text(19, 0, 'Deceased')])

```



```

# number of customers who churned and not churned
exit_counts = data['Churn Value'].value_counts()
exit_percentages = exit_counts

sns.set_style('whitegrid')
plt.figure(figsize=(8,6))
ax = sns.barplot(x=exit_counts.index, y=exit_counts.values, palette='pastel')
ax.set(xlabel='Churn Customers', ylabel='Number of Customers', title='Distribution of Churn Customers')
plt.xticks([0, 1], ['Not Churn', 'Churn'])
plt.ylim(top=max(exit_counts.values)*1.1)

# add counting number on top of each bar


```

```

for i, v in enumerate(exit_percentages):
    ax.text(i, exit_counts.values[i]+30, f'{v}', fontsize=14, color='black', ha='center')

plt.tight_layout()
plt.show()

```



```

# Does having partner or being senior citizen affect churns ?
fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(15,6))

```

```

colors = ['#1a237e', '#FFA500']
partner_churn = data.groupby(['Partner', 'Churn Value'])['Churn Value'].count().unstack().plot(ax=ax1, kind='bar', color=co
ax1.set_xlabel('Partner')
ax1.set_ylabel('Count')
ax1.set_title('Customer Churn by Partner')
ax1.legend(['Non-Churn', 'Churn'], title='Churn Value', loc='upper right')

```

```


senior_churn = data.groupby(['Senior Citizen', 'Churn Value'])['Churn Value'].count().unstack().plot(ax=ax2, kind='bar', co
ax2.set_xlabel('Senior Citizen')
ax2.set_ylabel('Count')
ax2.set_title('Customer Churn by Senior Citizen')
ax2.legend(['Non-Churn', 'Churn'], title='Churn Value', loc='upper right')

```

```

plt.suptitle('Customer Churn by Partner and Senior Citizen', fontsize=16, fontweight='bold')
plt.tight_layout()
plt.show()

```




```
# Contract wise churned or not ?
churn_labels = data['Churn Label'].unique()
contracts = data['Contract'].unique()

fig, ax = plt.subplots(figsize=(7, 5))

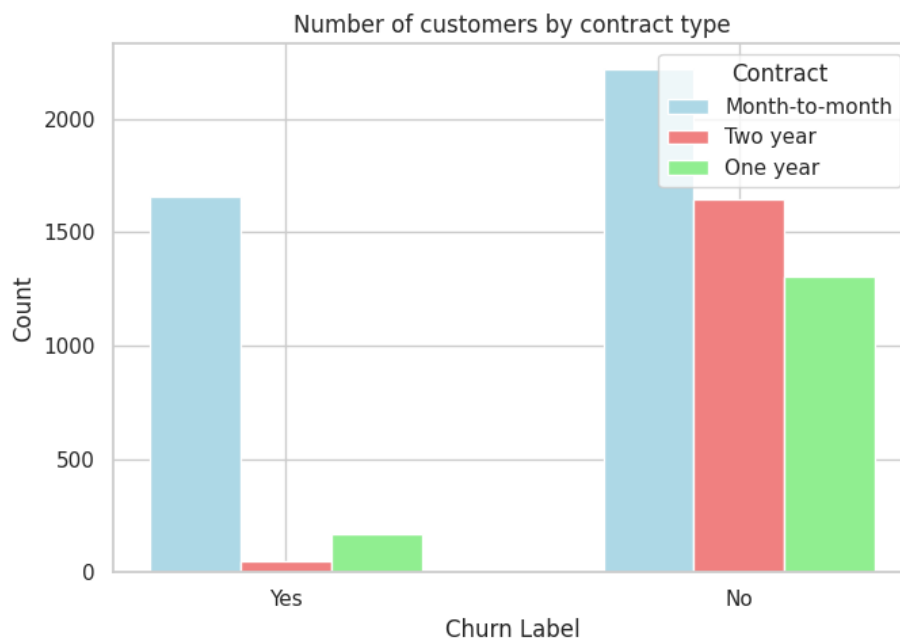
contract_colors = {'Month-to-month': 'lightblue', 'One year': 'lightgreen', 'Two year': 'lightcoral'}

width = 0.2
x = np.arange(len(churn_labels))

for i, contract in enumerate(contracts):
    contract_data = data[data['Contract'] == contract]
    counts = [contract_data[contract_data['Churn Label'] == label]['Churn Label'].count() for label in churn_labels]
    ax.bar(x + i * width, counts, width=width, label=contract, color=contract_colors[contract])

ax.set_xlabel('Churn Label')
ax.set_ylabel('Count')
ax.set_title('Number of customers by contract type')
ax.set_xticks(x + width)
ax.set_xticklabels(churn_labels)
ax.legend(title='Contract', loc='upper right')

plt.tight_layout()
plt.show()
```



```
#
grouped = data.groupby(['Contract', 'Churn Label'])['CustomerID'].count().reset_index()

fig = px.pie(data.groupby(['Contract', 'Churn Label'])['CustomerID'].count().reset_index(),
             values='CustomerID',
             names='Contract',
             facet_col = 'Churn Label',
             title = 'Churn rate by contract type')

fig.show()
```

Churn rate by contract type


```
Churn Label=Yes
# Does contract type affect Tenure months & churn ?
data.groupby(['Contract','Churn Label'])['Tenure Months'].mean()
```

Contract	Churn Label	
Month-to-month	No	21.054141
	Yes	14.350878
One year	No	41.674063
	Yes	44.963855
Two year	No	56.602914
	Yes	61.270833

Name: Tenure Months, dtype: float64

Total charges


```
plt.hist(data['Total Charges'], bins=50)
(array([1101., 547., 404., 337., 312., 270., 272., 270., 235.,
        180., 190., 135., 123., 121., 116., 104., 93., 115.,
        99., 92., 100., 79., 99., 100., 82., 80., 86.,
        86., 102., 60., 71., 75., 93., 74., 83., 62.,
        73., 62., 58., 56., 57., 44., 44., 44., 35.,
        41., 29., 26., 18., 8.]),
array([ 0., 173.696, 347.392, 521.088, 694.784, 868.48 ,
        1042.176, 1215.872, 1389.568, 1563.264, 1736.96 , 1910.656,
        2084.352, 2258.048, 2431.744, 2605.44 , 2779.136, 2952.832,
        3126.528, 3300.224, 3473.92 , 3647.616, 3821.312, 3995.008,
        4168.704, 4342.4 , 4516.096, 4689.792, 4863.488, 5037.184,
        5210.88 , 5384.576, 5558.272, 5731.968, 5905.664, 6079.36 ,
        6253.056, 6426.752, 6600.448, 6774.144, 6947.84 , 7121.536,
        7295.232, 7468.928, 7642.624, 7816.32 , 7990.016, 8163.712,
        8337.408, 8511.104, 8684.8 ]),
<BarContainer object of 50 artists>)
```



Monthly Charges

```
plt.hist(data['Monthly Charges'], bins=50)
```

```
(array([868., 316., 66., 343., 13., 59., 25., 25., 74., 8., 62.,
       32., 57., 194., 19., 190., 107., 96., 204., 43., 148., 85.,
       67., 125., 48., 278., 116., 166., 265., 66., 308., 156., 150.,
       259., 82., 270., 130., 165., 213., 93., 236., 114., 149., 177.,
       82., 115., 58., 50., 54., 17.]),
 array([ 18.25, 20.26, 22.27, 24.28, 26.29, 28.3 , 30.31, 32.32,
        34.33, 36.34, 38.35, 40.36, 42.37, 44.38, 46.39, 48.4 ,
        50.41, 52.42, 54.43, 56.44, 58.45, 60.46, 62.47, 64.48,
        66.49, 68.5 , 70.51, 72.52, 74.53, 76.54, 78.55, 80.56,
        82.57, 84.58, 86.59, 88.6 , 90.61, 92.62, 94.63, 96.64,
        98.65, 100.66, 102.67, 104.68, 106.69, 108.7 , 110.71, 112.72,
        114.73, 116.74, 118.75]),
 <BarContainer object of 50 artists>)
```



```
data.groupby('Churn Label')['Monthly Charges'].quantile([.50,.75,.95,.99])
```

```
Churn Label
No          0.50      64.4250
           0.75      88.4000
           0.95     108.4175
           0.99     115.1000
Yes         0.50      79.6500
           0.75      94.2000
           0.95     105.6100
           0.99     111.1320
Name: Monthly Charges, dtype: float64
```

```
corr_df = data.copy()
```

```
corr_df['Churn Label'].replace(to_replace='Yes', value=1, inplace=True)
corr_df['Churn Label'].replace(to_replace='No', value=0, inplace=True)
```


```
df_dummies = pd.get_dummies(corr_df[['Churn Label', 'Phone Service', 'Multiple Lines', 'Internet Service', 'Online Security',
                                       'Online Backup', 'Device Protection', 'Tech Support', 'Streaming TV',
                                       'Streaming Movies']])
df_dummies.head()
```

	Churn Label	Phone Service_No	Phone Service_Yes	Multiple Lines_No	Multiple Lines_Yes	Internet Service_DSL	Internet Service_Fiber optic	Internet Service_No	Online Security_No	...
0	1	0	1	1	0	0	1	0	0	0
1	1	0	1	1	0	0	0	1	0	1
2	1	0	1	0	0	1	0	1	0	1
3	1	0	1	0	0	1	0	1	0	1
4	1	0	1	0	0	1	0	1	0	1

```
5 rows x 27 columns
```

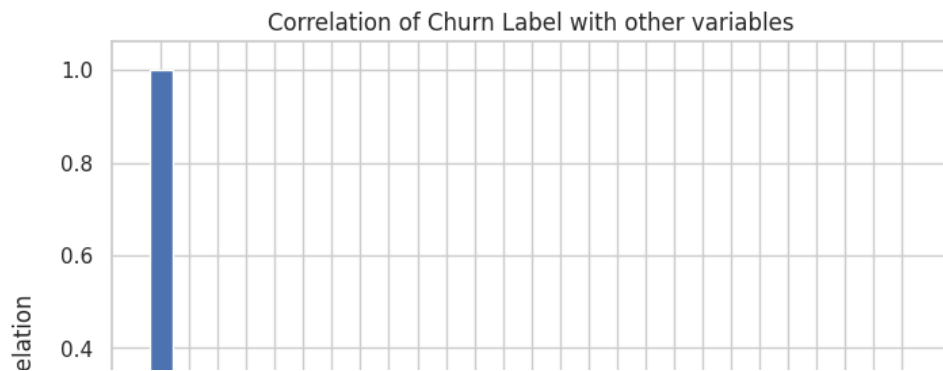
```
# Plot correlation of features
plt.figure(figsize=(9, 7))
sns.heatmap(df_dummies.corr(), annot=False, cmap='coolwarm')

plt.show()
```



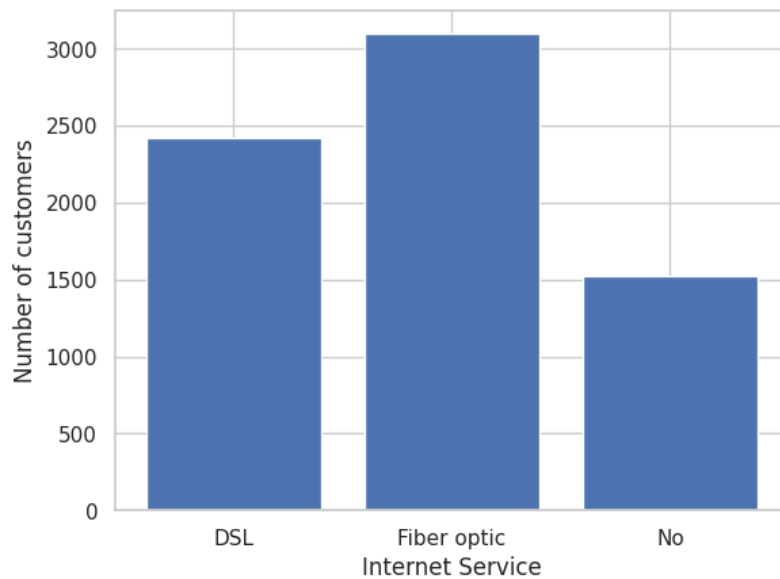
Correlation of Churn Label with other variables

```
fig, ax = plt.subplots(figsize=(8, 6))
ax.bar(df_dummies.corr()['Churn Label'].sort_values(ascending=False).index,
      df_dummies.corr()['Churn Label'].sort_values(ascending=False).values)
ax.set_title('Correlation of Churn Label with other variables')
ax.set_xlabel('Variables')
ax.set_ylabel('Correlation')
plt.xticks(rotation=90)
plt.show()
```




```
# Customers subscribed to each internet service
internet_services = data.groupby('Internet Service')['CustomerID'].count().reset_index()
plt.bar(internet_services['Internet Service'], internet_services['CustomerID'])
plt.xlabel('Internet Service')
plt.ylabel('Number of customers')
```

Text(0, 0.5, 'Number of customers')




```
data.groupby('Gender')['CustomerID'].count().plot(kind='pie')
plt.ylabel('')
```

Text(0, 0.5, '')




```
#
fig = px.pie(data.groupby(['Internet Service', 'Churn Label'])['CustomerID'].count().reset_index(),
             values='CustomerID',
             facet_col = 'Churn Label',
             names='Internet Service',
             title = 'What type of internet was connected to the clients who left the service?')
fig.show()
```

What type of internet was connecte



Tech Support

```
# Tech support opted ?
fig = px.bar(data.groupby(['Internet Service', 'Tech Support', 'Churn Label'])['CustomerID'].count().reset_index(),
             x="Internet Service",
             y="CustomerID",
             color="Churn Label",
             text = 'CustomerID',
             barmode="group",
             facet_col="Tech Support"
            )
fig.show()
```




```
# Assuming 'data' is a DataFrame with the required data
plt.figure(figsize=(12, 6)) # Adjust the figure size as needed

# Create a catplot with two pie charts for each 'Churn Label'
g = sns.catplot(data=data, x='Tech Support', kind='count', hue='Churn Label', aspect=2)
g.set_axis_labels('Tech Support', 'Count')
g.set_titles('Tech support option and churn ({col_name})')
```


```
plt.show()
```

<Figure size 1200x600 with 0 Axes>



```
fig = px.pie(data.groupby(['Tech Support', 'Churn Label'])['CustomerID'].count().reset_index(),
             values='CustomerID',
             facet_col = 'Churn Label',
             hole = .5,
             names='Tech Support',
             title = 'Tech support option and churn')
fig.show()
```


Tech support option and churn




```
internet_services = data.groupby('Internet Service')['CustomerID'].count().reset_index()
```

```
plt.bar(internet_services['Internet Service'], internet_services['CustomerID'])
```


<BarContainer object of 3 artists>



```
fig = px.bar(data.groupby(['Internet Service', 'Online Security',
                           'Churn Label'])['CustomerID'].count().reset_index(),
             x="Internet Service",
             y="CustomerID",
             color="Churn Label",
             #barmode="group",
             text = 'CustomerID',
             facet_col = 'Online Security'
            )
fig.show()
```



```
fig = px.bar(data.groupby(['Payment Method',
                           'Churn Label'])['CustomerID'].count().reset_index(),
             x="CustomerID",
             y="Payment Method",
             color="Churn Label",
             text = 'CustomerID'
            )
fig.show()
```

```
grouped = data.groupby(['Payment Method', 'Internet Service'])['CustomerID'].count().reset_index()
```

```
plt.figure(figsize=(12, 6))
```

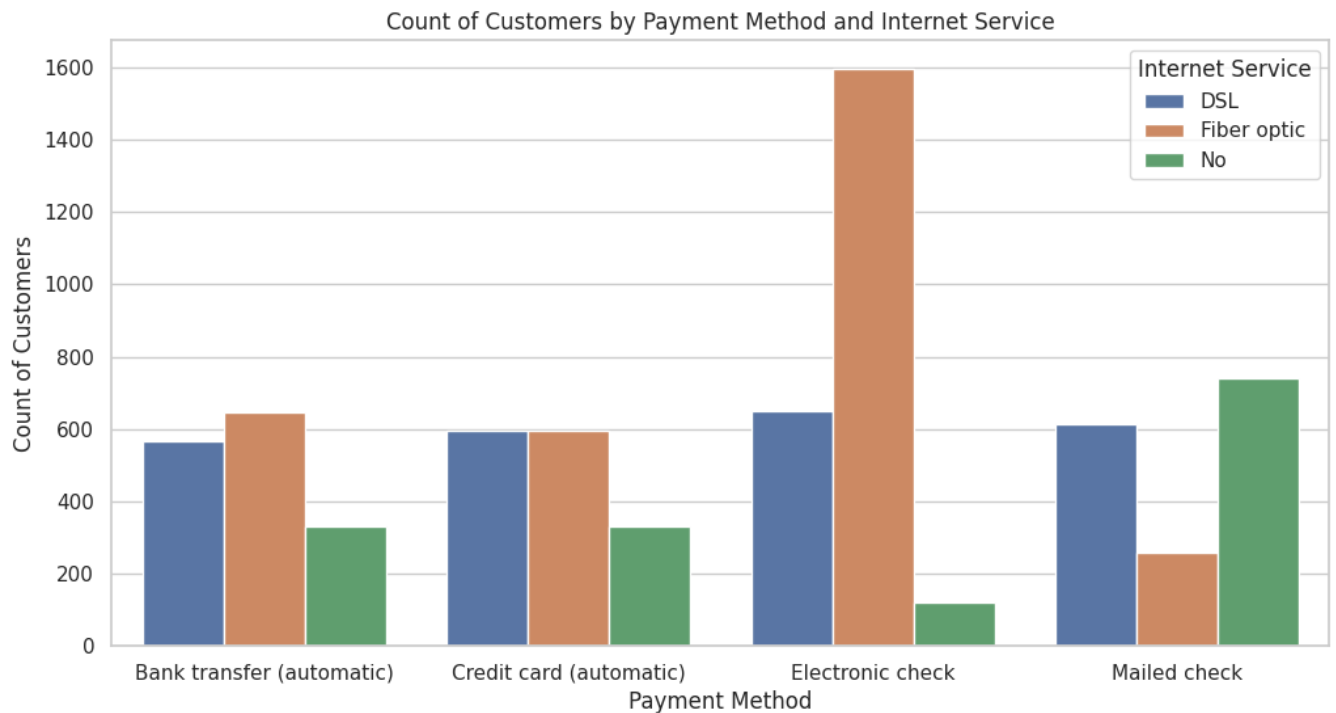
```
sns.barplot(data=grouped, x='Payment Method', y='CustomerID', hue='Internet Service')
```

```
plt.title("Count of Customers by Payment Method and Internet Service")
```

```
plt.xlabel("Payment Method")
```

```
plt.ylabel("Count of Customers")
```

```
plt.show()
```



```
churn_pm = data.assign(churn_clients = np.where(data['Churn Label'] == 'Yes', data['CustomerID'], None))\
    .groupby(['Payment Method', 'Internet Service']).agg({'churn_clients': 'count'}).reset_index()
```

```
pm_clients = data.groupby(['Payment Method', 'Internet Service'])['CustomerID'].count().reset_index()
```

```
pm_data = pm_clients.join(churn_pm.set_index(['Payment Method', 'Internet Service']), on=['Payment Method', 'Internet Service'])
```

```
pm_data
```

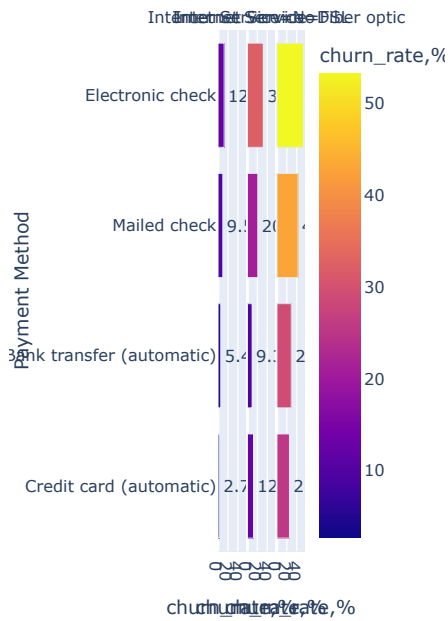
	Payment Method	Internet Service	CustomerID	churn_clients
0	Bank transfer (automatic)	DSL	566	53
1	Bank transfer (automatic)	Fiber optic	646	187

```

pm_data['churn_rate,%'] = round(((pm_data['churn_clients']/pm_data['CustomerID']) * 100),2)

# Churn rate % based on internet service & distributed based on payment method
fig = px.bar(pm_data.sort_values('churn_rate,%'),
             x='churn_rate,%',
             y='Payment Method',
             facet_col = 'Internet Service',
             color = 'churn_rate,%',
             text = 'churn_rate,%')
fig.show()

```



```


fig = px.bar(data.groupby(['Gender', 'Churn Label'])['CustomerID'].count().reset_index(),
             x="CustomerID",
             y="Gender",
             color="Churn Label",
             text = 'CustomerID'
             )
fig.show()

```

```
fig = px.pie(data.groupby(['Senior Citizen', 'Churn Label'])['CustomerID'].count().reset_index(),
             values='CustomerID',
             names='Churn Label',
             facet_col = 'Senior Citizen',
             color = 'Churn Label',
             title = 'Churn rate by customer age')

fig.show()
```

Churn rate by customer age




```
data.groupby('Senior Citizen')['CustomerID'].count()
```


```
Senior Citizen
No      5901
Yes     1142
Name: CustomerID, dtype: int64
```

```
fig = px.bar(data.groupby(['Senior Citizen', 'Partner',
                           'Dependents', 'Churn Label'])['CustomerID'].count().reset_index(),
             x="Senior Citizen",
             y="CustomerID",
             color="Churn Label",
             #barmode="group",
             facet_row="Partner",
             facet_col = 'Dependents'
             )

fig.show()
```



```
fig = px.bar(data.groupby(['Senior Citizen','Internet Service','Churn Label'])['CustomerID'].count().reset_index(),
             x="Internet Service",
             y="CustomerID",
             color="Churn Label",
             barmode="group",
             facet_col = 'Senior Citizen'
             )
fig.show()
```



```
# data.to_excel('/content/drive/MyDrive/datasets/fds_customer_churn/Telco_customer_churn_working.xlsx')
```

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.

▼ Name - Omkar Phansopkar

UID - 2022701007

D3 - CSE-DS

Dataset - Telco customer churn: IBM dataset

<https://www.kaggle.com/datasets/yeanzc/telco-customer-churn-ibm-dataset>

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
!pip install xgboost --quiet
```

```
import joblib
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib
import imblearn
import os
import folium
```

```
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder
```

```
from sklearn.linear_model import LogisticRegression
from xgboost import XGBClassifier
from sklearn.ensemble import RandomForestClassifier
from collections import Counter
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import accuracy_score, f1_score, recall_score, precision_score
from sklearn.metrics import average_precision_score, roc_auc_score, roc_curve, auc
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import RobustScaler
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
WORKING_DIR = '/content/drive/MyDrive/datasets/fds_customer_churn'
```

```
data = pd.read_excel(f'{WORKING_DIR}/Telco_customer_churn_working.xlsx', index_col=False)
data.sample(5)
```

	Unnamed: 0	CustomerID	Count	Country	State	City	Zip Code	Lat	Long
3511	3511	5242-UOWHD	1	United States	California	Anaheim	92802	33.807864,	-117.923782
1782	1782	6230-BSUXY	1	United States	California	Paramount	90723	33.897122,	-118.164432
2116	2116	9498-FIMXL	1	United States	California	Oceanside	92054	33.351059,	-117.420557
5479	5479	4393-RYCRE	1	United States	California	Macdoel	96058	41.769709,	-121.92063
2240	2240	0505-SPOOW	1	United States	California	Riverside	92508	33.885499,	-117.324959

5 rows x 35 columns

```
data = data.drop(['Unnamed: 0', 'Country', 'State', 'Count', 'Zip Code', 'Churn Reason', 'City', 'Churn Score', 'Churn Value', 'CLT Latitude', 'Longitude'], axis = 1)
```

```
data.sample(2)
```

	Gender	Senior Citizen	Partner	Dependents	Tenure Months	Phone Service	Multiple Lines	Internet Service
1598	Female	No	No	No	1	Yes	No	Fiber optic
4342	Male	No	Yes	Yes	23	Yes	No	No

2 rows x 21 columns

```
data.columns
```

```
Index(['Gender', 'Senior Citizen', 'Partner', 'Dependents', 'Tenure Months',  
      'Phone Service', 'Multiple Lines', 'Internet Service',  
      'Online Security', 'Online Backup', 'Device Protection', 'Tech Support',  
      'Streaming TV', 'Streaming Movies', 'Contract', 'Paperless Billing',  
      'Payment Method', 'Monthly Charges', 'Total Charges', 'Churn Label',  
      'hex_id'],  
      dtype='object')
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 7043 entries, 0 to 7042  
Data columns (total 21 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   Gender                7043 non-null  object  
1   Senior Citizen        7043 non-null  object  
2   Partner               7043 non-null  object  
3   Dependents            7043 non-null  object  
4   Tenure Months         7043 non-null  int64  
5   Phone Service         7043 non-null  object  
6   Multiple Lines        7043 non-null  object  
7   Internet Service      7043 non-null  object  
8   Online Security       7043 non-null  object  
9   Online Backup         7043 non-null  object  
10  Device Protection     7043 non-null  object  
11  Tech Support          7043 non-null  object  
12  Streaming TV          7043 non-null  object  
13  Streaming Movies      7043 non-null  object  
14  Contract              7043 non-null  object  
15  Paperless Billing      7043 non-null  object  
16  Payment Method        7043 non-null  object  
17  Monthly Charges       7043 non-null  float64  
18  Total Charges         7043 non-null  float64  
19  Churn Label           7043 non-null  object  
20  hex_id               7043 non-null  object  
dtypes: float64(2), int64(1), object(18)  
memory usage: 1.1+ MB
```



```
data['Churn Label'].replace(to_replace='Yes', value=1, inplace=True)  
data['Churn Label'].replace(to_replace='No', value=0, inplace=True)
```

```
data['Churn Label'].value_counts()
```

```
0    5174  
1    1869  
Name: Churn Label, dtype: int64
```

```
data.corr()
```

```
<ipython-input-43-c44ded798807>:1: FutureWarning: The default value of numeri  
data.corr()
```

	Tenure Months	Monthly Charges	Total Charges	Churn Label	
Tenure Months	1.000000	0.247900	0.826178	-0.352229	
Monthly Charges	0.247900	1.000000	0.651174	0.193356	
Total Charges	0.826178	0.651174	1.000000	-0.198324	
Churn Label	-0.352229	0.193356	-0.198324	1.000000	

```

corr = data.corr()['Churn Label'].sort_values(ascending=False)

plt.figure(figsize=(10, 6))
bars = plt.bar(corr.index, corr, color=np.where(corr > 0, 'b', 'r'), alpha=0.7)

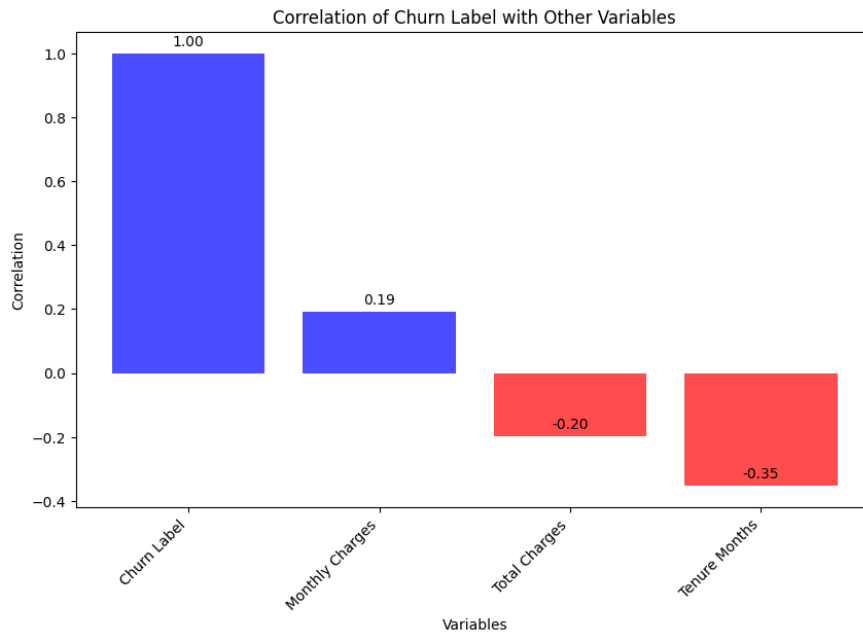
plt.title('Correlation of Churn Label with Other Variables')
plt.xlabel('Variables')
plt.ylabel('Correlation')
plt.xticks(rotation=45, ha='right')

for bar in bars:
    height = bar.get_height()
    plt.annotate(f'{height:.2f}', xy=(bar.get_x() + bar.get_width() / 2, height),
                xytext=(0, 3), textcoords='offset points', ha='center', va='bottom')

plt.show()

```

<ipython-input-44-b1b0091519ea>:1: FutureWarning: The default value of numeri
corr = data.corr()['Churn Label'].sort_values(ascending=False)



```

# Columns to apply nominal encoding
categorical_columns = []
for col in data.columns:
    if data[col].nunique() < 20:
        categorical_columns.append(col)
        print(col, "->", data[col].unique())

categorical_columns.remove("Churn Label")
categorical_columns

Gender -> ['Male' 'Female']
Senior Citizen -> ['No' 'Yes']
Partner -> ['No' 'Yes']
Dependents -> ['No' 'Yes']
Phone Service -> ['Yes' 'No']
Multiple Lines -> ['No' 'Yes' 'No phone service']
Internet Service -> ['DSL' 'Fiber optic' 'No']
Online Security -> ['Yes' 'No' 'No internet service']
Online Backup -> ['Yes' 'No' 'No internet service']
Device Protection -> ['No' 'Yes' 'No internet service']
Tech Support -> ['No' 'Yes' 'No internet service']
Streaming TV -> ['No' 'Yes' 'No internet service']
Streaming Movies -> ['No' 'Yes' 'No internet service']
Contract -> ['Month-to-month' 'Two year' 'One year']
Paperless Billing -> ['Yes' 'No']
Payment Method -> ['Mailed check' 'Electronic check' 'Bank transfer (automatic)'
                  'Credit card (automatic)']
Churn Label -> [1 0]

```

```
['Gender',
 'Senior Citizen',
 'Partner',
 'Dependents',
 'Phone Service',
 'Multiple Lines',
 'Internet Service',
 'Online Security',
 'Online Backup',
 'Device Protection',
 'Tech Support',
 'Streaming TV',
 'Streaming Movies',
 'Contract',
 'Paperless Billing',
 'Payment Method']
```

data.columns

```
Index(['Gender', 'Senior Citizen', 'Partner', 'Dependents', 'Tenure Months',
       'Phone Service', 'Multiple Lines', 'Internet Service',
       'Online Security', 'Online Backup', 'Device Protection', 'Tech Support',
       'Streaming TV', 'Streaming Movies', 'Contract', 'Paperless Billing',
       'Payment Method', 'Monthly Charges', 'Total Charges', 'Churn Label',
       'hex_id'],
      dtype='object')
```

data.head(2)

	Gender	Senior Citizen	Partner	Dependents	Tenure Months	Phone Service	Multiple Lines	Internet Service	S
0	Male	No	No	No	2	Yes	No	DSL	
1	Female	No	No	Yes	2	Yes	No	Fiber optic	

2 rows × 21 columns

data.columns

```
Index(['Gender', 'Senior Citizen', 'Partner', 'Dependents', 'Tenure Months',
       'Phone Service', 'Multiple Lines', 'Internet Service',
       'Online Security', 'Online Backup', 'Device Protection', 'Tech Support',
       'Streaming TV', 'Streaming Movies', 'Contract', 'Paperless Billing',
       'Payment Method', 'Monthly Charges', 'Total Charges', 'Churn Label',
       'hex_id'],
      dtype='object')
```

```
def encode_data(df):
    # Nominal encoding
    encoder = OneHotEncoder(sparse_output=False, drop='first')

    encoded_data = encoder.fit_transform(data[categorical_columns])
    encoded_data = pd.DataFrame(encoded_data, columns=encoder.get_feature_names_out(categorical_columns))

    newdf = pd.concat([df, encoded_data], axis=1)
    return newdf.drop(columns=[*categorical_columns, 'hex_id'])
```

```
data = encode_data(data)
data.sample(3)
```

	Tenure Months	Monthly Charges	Total Charges	Churn Label	Gender_Male	Senior Citizen_Yes	Partner_Yes	D
3766	71	99.00	7061.65	0	0.0	0.0	1.0	
1199	27	85.25	2287.25	1	0.0	1.0	0.0	
3544	1	49.85	49.85	0	1.0	0.0	0.0	

3 rows × 31 columns

```
data.groupby('Churn Label')['Churn Label'].count()
```

Churn Label	
0	5174


```
1 1869
Name: Churn Label, dtype: int64
```

```
len(data.columns)
```

```
31
```

```
/usr/local/lib/python3.10/dist-packages/statsmodels/stats/outliers_influence.py:198: RuntimeWarning: divide by zero e
vif = 1. / (1. - r_squared_i)
```

	Feature	VIF
0	Tenure Months	7.542697
1	Monthly Charges	865.450546
2	Total Charges	10.862958
3	Churn Label	1.429089
4	Gender_Male	1.002024
5	Senior Citizen_Yes	1.135728
6	Partner_Yes	1.344290
7	Dependents_Yes	1.273662
8	Phone Service_Yes	1771.930567
9	Multiple Lines_No phone service	60.982066
10	Multiple Lines_Yes	7.281349
11	Internet Service_Fiber optic	148.398568
12	Internet Service_No	inf
13	Online Security_No internet service	inf
14	Online Security_Yes	6.339317
15	Online Backup_No internet service	inf
16	Online Backup_Yes	6.783802
17	Device Protection_No internet service	inf
18	Device Protection_Yes	6.927119
19	Tech Support_No internet service	inf
20	Tech Support_Yes	6.473185
21	Streaming TV_No internet service	inf
22	Streaming TV_Yes	24.073788
23	Streaming Movies_No internet service	inf
24	Streaming Movies_Yes	24.131570
25	Contract_One year	1.632110
26	Contract_Two year	2.629070
27	Paperless Billing_Yes	1.212537
28	Payment Method_Credit card (automatic)	1.560576
29	Payment Method_Electronic check	1.983301
30	Payment Method_Mailed check	1.860244
11		

```
over = SMOTE(sampling_strategy = 1)
```

```
x = data.drop("Churn Label", axis = 1).values
y = data['Churn Label'].values
```

```
x,y = over.fit_resample(x,y)
```

```
x,y
```

```
(array([[2.00000000e+00, 5.38500000e+01, 1.08150000e+02, ...,
0.00000000e+00, 0.00000000e+00, 1.00000000e+00],
[2.00000000e+00, 7.07000000e+01, 1.51650000e+02, ...,
0.00000000e+00, 1.00000000e+00, 0.00000000e+00],
[8.00000000e+00, 9.96500000e+01, 8.20500000e+02, ...,
0.00000000e+00, 1.00000000e+00, 0.00000000e+00],
...,
[2.00000000e+00, 8.72689071e+01, 1.83367799e+02, ...,
0.00000000e+00, 8.49336565e-02, 0.00000000e+00],
[1.00000000e+00, 4.25625192e+01, 4.25625192e+01, ...,
0.00000000e+00, 1.00000000e+00, 0.00000000e+00],
[4.62333136e+01, 1.07678530e+02, 4.87392507e+03, ...,
0.00000000e+00, 1.00000000e+00, 0.00000000e+00]]),
array([1, 1, 1, ..., 1, 1, 1]))
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, random_state=2, test_size = 0.2)
```

```
x_test
```

```
array([[1.70000000e+01, 8.82500000e+01, 1.46065000e+03, ...,
0.00000000e+00, 1.00000000e+00, 0.00000000e+00],
[8.00000000e+00, 5.13000000e+01, 4.11600000e+02, ...,
1.00000000e+00, 0.00000000e+00, 0.00000000e+00],
[1.00000000e+00, 7.11070051e+01, 7.11070051e+01, ...,
0.00000000e+00, 1.00000000e+00, 0.00000000e+00],
...,
[6.90000000e+01, 4.62500000e+01, 3.12140000e+03, ...,
0.00000000e+00, 0.00000000e+00, 0.00000000e+00],
[1.70000000e+01, 5.46000000e+01, 9.34800000e+02, ...,
```

```
0.00000000e+00, 1.00000000e+00, 0.00000000e+00],
[1.00000000e+00, 7.47689188e+01, 7.47689188e+01, ...,
0.00000000e+00, 1.00000000e+00, 0.00000000e+00]])
```

```
x_train
```

```
array([[3.00000000e+00, 2.05500000e+01, 5.74000000e+01, ...,
1.00000000e+00, 0.00000000e+00, 0.00000000e+00],
[1.00000000e+00, 4.02000000e+01, 4.02000000e+01, ...,
0.00000000e+00, 1.00000000e+00, 0.00000000e+00],
[1.50000000e+01, 2.54000000e+01, 3.99600000e+02, ...,
1.00000000e+00, 0.00000000e+00, 0.00000000e+00],
...,
[6.00000000e+00, 1.95500000e+01, 1.24450000e+02, ...,
0.00000000e+00, 0.00000000e+00, 1.00000000e+00],
[5.40000000e+01, 1.04100000e+02, 5.64580000e+03, ...,
0.00000000e+00, 0.00000000e+00, 0.00000000e+00],
[1.00000000e+00, 2.07741153e+01, 2.07741153e+01, ...,
0.00000000e+00, 0.00000000e+00, 1.00000000e+00]])
```

```
x_train[0]
```

```
array([ 3. , 20.55, 57.4 ,  0. ,  0. ,  0. ,  1. ,  1. ,  0. ,
        0. ,  0. ,  1. ,  1. ,  0. ,  1. ,  0. ,  1. ,  0. ,
        1. ,  0. ,  1. ,  0. ,  1. ,  0. ,  1. ,  0. ,  1. ,
        1. ,  0. ,  0.  ])
```

```
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
# Filter only the numerical columns
```

```
numerical_data = data.select_dtypes(include=['number'])
```

```
# VIF for each feature
```

```
vif = pd.DataFrame()
```

```
vif["Feature"] = numerical_data.columns
```

```
vif["VIF"] = [variance_inflation_factor(numerical_data.values, i) for i in range(numerical_data.shape[1])]
```

```
# Select features with VIF below threshold 6
```

```
remove_features = vif[vif["VIF"] < 6]["Feature"].values
```

```
remove_features
```

```
array(['Monthly Charges', 'Churn Label'], dtype=object)
```

```
data.drop(columns = remove_features)
```

	Gender	Senior Citizen	Partner	Dependents	Tenure Months	Phone Service	Multiple Lines	Internet Service
0	Male	No	No	No	2	Yes	No	DSL

```

def test_model(model, x_train, y_train, x_test, y_test):
    model.fit(x_train, y_train)

    scaler = StandardScaler()
    x_train = scaler.fit_transform(x_train)

    predictions = model.predict(x_test)
    c_matrix = confusion_matrix(y_test, predictions)

    percentages = (c_matrix / np.sum(c_matrix, axis=1)[:, np.newaxis]).round(2) * 100
    labels = [[f'{c_matrix[i, j]} ({percentages[i, j]:.2f}%)' for j in range(c_matrix.shape[1])] for i in range(c_matrix.shape[0])]
    labels = np.asarray(labels)

    sns.heatmap(c_matrix, annot=labels, fmt='', cmap='Blues')

    print("ROC AUC: ", '{:.2%}'.format(roc_auc_score(y_test, predictions)))
    print("Model accuracy: ", '{:.2%}'.format(accuracy_score(y_test, predictions)))
    print(classification_report(y_test, predictions))

def predict_sample(model, sample):
    prediction = model.predict([sample])[0]
    scaler = StandardScaler()
    sample = scaler.fit_transform([sample])
    prediction = model.predict(sample)[0]
    return prediction

def save_model(model, filename):
    try:
        joblib.dump(model, filename)
        print(f"Model saved to {filename}")
    except Exception as e:
        print(f"Error saving the model: {str(e)}")

# logistic_reg = LogisticRegression()
# test_model(logistic_reg, x_train, y_train, x_test, y_test)
# print(predict_sample(logistic_reg, x_test[0]))
# print(predict_sample(logistic_reg, x_test[1]))

# Train & test different models
def test_model(model, x_train, y_train, x_test, y_test):
    model.fit(x_train, y_train)

    predictions = model.predict(x_test)
    c_matrix = confusion_matrix(y_test, predictions)

    percentages = (c_matrix / np.sum(c_matrix, axis=1)[:, np.newaxis]).round(2) * 100
    labels = [[f'{c_matrix[i, j]} ({percentages[i, j]:.2f}%)' for j in range(c_matrix.shape[1])] for i in range(c_matrix.shape[0])]
    labels = np.asarray(labels)

    sns.heatmap(c_matrix, annot=labels, fmt='', cmap='Blues')

    print("ROC AUC: ", '{:.2%}'.format(roc_auc_score(y_test, predictions)))
    print("Model accuracy: ", '{:.2%}'.format(accuracy_score(y_test, predictions)))

    mae = mean_absolute_error(y_test, predictions)
    rmse = np.sqrt(mean_squared_error(y_test, predictions))
    print(f'Mean Absolute Error (MAE): {mae}')
    print(f'Mean squared error (MSE): {mean_squared_error(y_test, predictions)}')
    print(f'Root Mean Squared Error (RMSE): {rmse}')
    print(f'Sum squared error (SSE): {np.sum((y_test - predictions) ** 2)}')
    print(f'Sum squared total (SST): {np.sum((y_test - np.mean(y_test)) ** 2)}')
    print(f'R^2: {r2_score(y_test, predictions)}')
    print("\n\n")
    print(classification_report(y_test, predictions))

# Predict single sample
def predict_sample(model, sample):
    prediction = model.predict([sample])[0]
    return prediction

# Save pickle
def save_model(model, filename):
    try:
        joblib.dump(model, filename)

```

```

print(f"Model saved to {filename}")
except Exception as e:
    print(f"Error saving the model: {str(e)}")

logistic_reg = LogisticRegression()
test_model(logistic_reg, x_train, y_train, x_test, y_test)
save_model(logistic_reg, f"{WORKING_DIR}/logistic_reg.pkl")


/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regres
n_iter_i = _check_optimize_result(
ROC AUC: 78.47%
Model accuracy: 78.41%
Mean Absolute Error (MAE): 0.21594202898550724
Mean squared error (MSE) 0.21594202898550724
Root Mean Squared Error (RMSE): 0.4646956304781736
Sum squared error (SSE): 447
Sum squared total (SST): 517.4053140096617
R^2: 0.13607381312737532

```

	precision	recall	f1-score	support
0	0.82	0.73	0.78	1049
1	0.75	0.84	0.79	1021
accuracy			0.78	2070
macro avg	0.79	0.78	0.78	2070
weighted avg	0.79	0.78	0.78	2070

Model saved to /content/drive/MyDrive/datasets/fds_customer_churn/logistic_re



```


logistic_reg_high = LogisticRegression(max_iter=4000)
test_model(logistic_reg_high, x_train, y_train, x_test, y_test)
save_model(logistic_reg_high, f"{WORKING_DIR}/logistic_reg_high.pkl")

```

ROC AUC: 78.24%
 Model accuracy: 78.16%
 Mean Absolute Error (MAE): 0.21835748792270532
 Mean squared error (MSE) 0.21835748792270532
 Root Mean Squared Error (RMSE): 0.4672873718844811
 Sum squared error (SSE): 452
 Sum squared total (SST): 517.4053140096617
 R²: 0.12641020924736834

	precision	recall	f1-score	support
0	0.82	0.73	0.77	1049
1	0.75	0.84	0.79	1021
accuracy			0.78	2070
macro avg	0.79	0.78	0.78	2070
weighted avg	0.79	0.78	0.78	2070

Model saved to /content/drive/MyDrive/datasets/fds_customer_churn/logistic_re




```
xgb = XGBClassifier(learning_rate= 0.01,max_depth = 3,n_estimators = 1000)
test_model(xgb,x_train,y_train,x_test,y_test)
save_model(xgb, f"{WORKING_DIR}/xgb.pkl")
```

ROC AUC: 86.25%
 Model accuracy: 86.23%
 Mean Absolute Error (MAE): 0.13768115942028986
 Mean squared error (MSE) 0.13768115942028986
 Root Mean Squared Error (RMSE): 0.3710541192606408
 Sum squared error (SSE): 285
 Sum squared total (SST): 517.4053140096617
 R²: 0.44917457883960177

	precision	recall	f1-score	support
0	0.88	0.85	0.86	1049
1	0.85	0.88	0.86	1021
accuracy			0.86	2070
macro avg	0.86	0.86	0.86	2070
weighted avg	0.86	0.86	0.86	2070

Model saved to /content/drive/MyDrive/datasets/fds_customer_churn/xgb.pkl




```
rf_classifier = RandomForestClassifier(n_estimators=500, random_state=42)
test_model(rf_classifier,x_train,y_train,x_test,y_test)
save_model(rf_classifier, f"{WORKING_DIR}/rf_classifier.pkl")
```

ROC AUC: 86.12%
 Model accuracy: 86.14%
 Mean Absolute Error (MAE): 0.1386473429951691
 Mean squared error (MSE) 0.1386473429951691
 Root Mean Squared Error (RMSE): 0.37235378740543124
 Sum squared error (SSE): 287
 Sum squared total (SST): 517.4053140096617
 R²: 0.4453091372875989

	precision	recall	f1-score	support
0	0.86	0.87	0.86	1049
1	0.87	0.85	0.86	1021
accuracy			0.86	2070
macro avg	0.86	0.86	0.86	2070
weighted avg	0.86	0.86	0.86	2070

Model saved to /content/drive/MyDrive/datasets/fds_customer_churn/rf_classifi



y_test[[3,9,20]]

array([1, 0, 1])

```

print(predict_sample(logistic_reg_high, x_test[3]))
print(predict_sample(logistic_reg_high, x_test[9]))
print(predict_sample(logistic_reg_high, x_test[20]))

```

1
0
1

```

print(predict_sample(xgb, x_test[3]))
print(predict_sample(xgb, x_test[9]))
print(predict_sample(xgb, x_test[20]))

```

1
0
1

```

print(predict_sample(rf_classifier, x_test[3]))
print(predict_sample(rf_classifier, x_test[9]))
print(predict_sample(rf_classifier, x_test[20]))

```

1
0
1

