

Get data set from Kaggel winemag-data-130k-v2.csv 33608 entries, 0 to 33607 Data columns (total 13 columns)

```
import pandas as pd
import numpy as np

# reviews = pd.read_csv("winemag-data-130k-v2.csv", index_col=0)
reviews = pd.read_csv("winemag-data-130k-v2_full.csv", index_col=0)

reviews.head()
```

	country	description	designation	points	price	province	region_1	region_2
0	Italy	Aromas include tropical fruit, broom, brimston...	Vulkà Bianco	87	NaN	Sicily & Sardinia	Etna	
1	Portugal	This is ripe and fruity, a wine that is smooth...	Avidagos	87	15.0	Douro	NaN	
2	US	Tart and snappy, the flavors of lime flesh and...	NaN	87	14.0	Oregon	Willamette Valley	Wi
3	US	Pineapple rind, lemon pith and orange blossom ...	Reserve Late Harvest	87	13.0	Michigan	Lake Michigan Shore	
4	US	Much like the regular bottling from 2012, this...	Vintner's Reserve Wild Child Block	87	65.0	Oregon	Willamette Valley	Wi

```
len(reviews)

129971
```

# rename column region\_1 as region and region\_2 as locale

```
renamedReviews = reviews.rename(columns = {'region_1':'region', 'region_2':'locale'})
renamedReviews.head()
```

	country	description	designation	points	price	province	region	
0	Italy	Aromas include tropical fruit, broom, brimston...	Vulkà Bianco	87	NaN	Sicily & Sardinia	Etna	
1	Portugal	This is ripe and fruity, a wine that is smooth...	Avidagos	87	15.0	Douro	NaN	
2	US	Tart and snappy, the flavors of lime flesh and...	NaN	87	14.0	Oregon	Willamette Valley	Wi
3	US	Pineapple rind, lemon pith and orange blossom ...	Reserve Late Harvest	87	13.0	Michigan	Lake Michigan Shore	
4	US	Much like the regular bottling from 2012, this...	Vintner's Reserve Wild Child Block	87	65.0	Oregon	Willamette Valley	Wi

get info of dataframe

```
reviews.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 129971 entries, 0 to 129970
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  -
0   country              129908 non-null object
```

```

1  description          129971 non-null object
2  designation          92506 non-null object
3  points              129971 non-null int64
4  price               120975 non-null float64
5  province            129908 non-null object
6  region_1            108724 non-null object
7  region_2            50511 non-null object
8  taster_name         103727 non-null object
9  taster_twitter_handle 98758 non-null object
10 title               129971 non-null object
11 variety             129970 non-null object
12 winery              129971 non-null object
dtypes: float64(1), int64(1), object(11)
memory usage: 13.9+ MB

```

```
renamedReviews.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 129971 entries, 0 to 129970
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   country               129908 non-null object
1   description           129971 non-null object
2   designation           92506 non-null object
3   points               129971 non-null int64
4   price                120975 non-null float64
5   province              129908 non-null object
6   region               108724 non-null object
7   locale               50511 non-null object
8   taster_name          103727 non-null object
9   taster_twitter_handle 98758 non-null object
10  title                129971 non-null object
11  variety              129970 non-null object
12  winery               129971 non-null object
dtypes: float64(1), int64(1), object(11)
memory usage: 13.9+ MB

```

Create a variable `df` containing the `country`, `province`, `region_1`, and `region_2` columns of the records with the index labels 0, 1, 10, and 100

```
df = reviews[['country', 'province', 'region_1', 'region_2']]
df.head()
```

	country	province	region_1	region_2
0	Italy	Sicily & Sardinia	Etna	NaN
1	Portugal	Douro	NaN	NaN
2	US	Oregon	Willamette Valley	Willamette Valley
3	US	Michigan	Lake Michigan Shore	NaN
4	US	Oregon	Willamette Valley	Willamette Valley

```
df.iloc[[0,1,10,100]]
```

	country	province	region_1	region_2	
<b>0</b>	Italy	Sicily & Sardinia	Etna	NaN	
<b>1</b>	Portugal	Douro	NaN	NaN	
<b>10</b>	US	California	Napa Valley	Napa	
<b>100</b>	US	New York	Finger Lakes	Finger Lakes	

What countries are represented in the review dataset? (Your answer should not include any duplicates.)

```
reviews['country'].unique()
```

```
array(['Italy', 'Portugal', 'US', 'Spain', 'France', 'Germany',
      'Argentina', 'Chile', 'Australia', 'Austria', 'South Africa',
      'New Zealand', 'Israel', 'Hungary', 'Greece', 'Romania', 'Mexico',
      'Canada', nan, 'Turkey', 'Czech Republic', 'Slovenia',
      'Luxembourg', 'Croatia', 'Georgia', 'Uruguay', 'England',
      'Lebanon', 'Serbia', 'Brazil', 'Moldova', 'Morocco', 'Peru',
      'India', 'Bulgaria', 'Cyprus', 'Armenia', 'Switzerland',
      'Bosnia and Herzegovina', 'Ukraine', 'Slovakia', 'Macedonia',
      'China', 'Egypt'], dtype=object)
```

How often does each country appear in the dataset? Create a Series `reviews_per_country` mapping countries to the count of reviews of medicines from that country.

```
reviews_per_country = reviews['country'].value_counts()
reviews_per_country
```

```
US          54504
France      22093
Italy       19540
Spain       6645
Portugal    5691
Chile       4472
Argentina   3800
Austria     3345
Australia   2329
Germany     2165
New Zealand 1419
South Africa 1401
Israel       505
Greece       466
Canada       257
Hungary      146
Bulgaria     141
Romania      120
Uruguay      109
```

Turkey	90
Slovenia	87
Georgia	86
England	74
Croatia	73
Mexico	70
Moldova	59
Brazil	52
Lebanon	35
Morocco	28
Peru	16
Ukraine	14
Serbia	12
Czech Republic	12
Macedonia	12
Cyprus	11
India	9
Switzerland	7
Luxembourg	6
Bosnia and Herzegovina	2
Armenia	2
Slovakia	1
China	1
Egypt	1

Name: country, dtype: int64

Create variable `centered_price` containing a version of the `price` column with the mean price subtracted.

(Note: this 'centering' transformation is a common preprocessing step before applying various machine learning algorithms.)

```
reviews['price'].mean()
```

```
35.363389129985535
```

```
centered_price = reviews['price'] - reviews['price'].mean()
centered_price
```

```
0      NaN
1    -20.363389
2    -21.363389
3    -22.363389
4     29.636611
...
129966  -7.363389
129967   39.636611
129968  -5.363389
129969  -3.363389
129970 -14.363389
Name: price, Length: 129971, dtype: float64
```

```
reviews.head()
```

	country	description	designation	points	price	province	region_1	r
0	Italy	Aromas include tropical fruit, broom, brimston...	Vulkà Bianco	87	NaN	Sicily & Sardinia	Etna	
1	Portugal	This is ripe and fruity, a wine that is smooth...	Avidagos	87	15.0	Douro	NaN	
2	US	Tart and snappy, the flavors of lime flesh and...	NaN	87	14.0	Oregon	Willamette Valley	Wi
3	US	Pineapple rind, lemon pith and orange blossom ...	Reserve Late Harvest	87	13.0	Michigan	Lake Michigan Shore	
4	US	Much like the regular bottling from 2012, this...	Vintner's Reserve Wild Child Block	87	65.0	Oregon	Willamette Valley	Wi

I'm an economical medicine buyer. Which medicine is the "best bargain"? Create a variable `bargain_medicine` with the title of the medicine with the highest points-to-price ratio in the dataset.

```
reviews.iloc[(reviews['points'] / reviews['price']).idxmax()]

country
description
designation
points
price
province
region_1
region_2
taster_name
taster_twitter_handle
title
variety
winery
```

```
US
There's a lot going on in this Merlot, which i...
NaN
86
4.0
California
California
California Other
NaN
NaN
Bandit NV Merlot (California)
Merlot
Bandit
```

```
Name: 64590, dtype: object
```

```
bargain_medicine = reviews.iloc[(reviews['points'] / reviews['price']).idxmax()]  
bargain_medicine
```

```
'Bandit NV Merlot (California)'
```

There are only so many words you can use when describing a bottle of medicine. Is a medicine more likely to be "tropical" or "fruity"? Create a Series descriptor\_counts counting how many times each of these two words appears in the description column in the dataset. (For simplicity, let's ignore the capitalized versions of

```
reviews['description'].map(lambda desc: 'tropical' in desc).sum()
```

```
3607
```

```
pd.Series({  
    'tropical': reviews['description'].map(lambda desc: 'tropical' in desc).sum(),  
    'fruity': reviews['description'].map(lambda desc: 'fruity' in desc).sum(),  
})
```

```
tropical    3607  
fruity      9090  
dtype: int64
```

We'd like to host these medicine reviews on our website, but a rating system ranging from 80 to 100 points is too hard to understand - we'd like to translate them into simple star ratings. A score of 95 or higher counts as 3 stars, a score of at least 85 but less than 95 is 2 stars. Any other score is 1 star.

Also, the Canadian Vintners Association bought a lot of ads on the site, so any medicines from Canada should automatically get 3 stars, regardless of points.

Create a series star\_ratings with the number of stars corresponding to each review in the dataset.

```
def pointMapper(row):  
    point = row.points  
    if row.country == 'Canada':  
        return 3  
    if point >= 95:  
        return 3  
    if point >= 85:  
        return 2  
    return 1  
star_ratings = reviews.apply(pointMapper, axis='columns')  
star_ratings
```

```

0      2
1      2
2      2
3      2
4      2
..
129966  2
129967  2
129968  2
129969  2
129970  2
Length: 129971, dtype: int64

```

1. What is the data type of the points column in the dataset?

```

reviews.points.dtype

dtype('int64')

```

3. Sometimes the price column is null. How many reviews in the dataset are missing a price?

```

reviews['price'].isnull().sum()

8996

```

4. What are the most common medicine-producing regions? Create a Series counting the number of times each value occurs in the region\_1 field. This field is often missing data, so replace missing values with Unknown. Sort in descending order. Your output should look something like this:

Unknown 21247

Napa Valley 4480

...

Bardolino Superiore 1

Primitivo del Tarantino 1

Name: region\_1, Length: 1230, dtype: int64

```

reviews['region_1'].fillna('Unknown', inplace=True)
reviews['region_1'].value_counts().sort_values(ascending=False)

```

```

Unknown      21247
Napa Valley   4480
Columbia Valley (WA)  4124
Russian River Valley  3091
California    2620

```



```

California      2023
...
Offida Rosso    1
Corton Perrières 1
Isle St. George 1
Geelong         1
Paestum         1
Name: region_1, Length: 1230, dtype: int64

```

Double-click (or enter) to edit

2. Create a Series from entries in the points column, but convert the entries to strings. Hint: strings are str in native Python.

```

reviews['points'].astype('str')

0      87
1      87
2      87
3      87
4      87
...
129966  90
129967  90
129968  90
129969  90
129970  90
Name: points, Length: 129971, dtype: object

```

Who are the most common medicine reviewers in the dataset? Create a Series whose index is the taster\_twitter\_handle category from the dataset, and whose values count how many reviews each person wrote.

```

reviews.groupby('taster_twitter_handle').size()

taster_twitter_handle
@AnneInVino      3685
@JoeCz           5147
@bkfiona         27
@gordone_cellars 4177
@kerinokeefe     10776
@laurbuzz        1835
@mattkettmann    6332
@paulgwine       9532
@suskostrzewa    1085
@vboone          9537
@vossroger       25514
@wawinereport    4966
@wineschach      15134
@winewchristina   6
@worldwineguys   1005
dtype: int64

```

2. What is the best medicine I can buy for a given amount of money? Create a Series whose index is medicine prices and whose values is the maximum number of points a medicine costing that much was given in a review. Sort the values by price, ascending (so that 4.0 dollars is at the top and 3300.0 dollars is at the bottom).

```
reviews.groupby('price')['points'].max()
```

```
price
4.0      86
5.0      87
6.0      88
7.0      91
8.0      91
..
1900.0    98
2000.0    97
2013.0    91
2500.0    96
3300.0    88
Name: points, Length: 390, dtype: int64
```

What are the minimum and maximum prices for each variety of medicine? Create a DataFrame whose index is the variety category from the dataset and whose values are the min and max values thereof.

```
minmax = reviews.groupby('variety').price.aggregate([min, max])
minmax
```

	min	max
variety		
<b>Abouriou</b>	15.0	75.0
<b>Agiorgitiko</b>	10.0	66.0
<b>Aglianico</b>	6.0	180.0
<b>Aidani</b>	27.0	27.0
<b>Airen</b>	8.0	10.0
...	...	...
<b>Zinfandel</b>	5.0	100.0
<b>Zlahtina</b>	13.0	16.0
<b>Zweigelt</b>	9.0	70.0
<b>Çalkarası</b>	19.0	19.0
<b>Žilavka</b>	15.0	15.0

707 rows × 2 columns

4. What are the most expensive medicine varieties? Create a variable `sorted_varieties` containing a copy of the dataframe from the previous question where varieties are sorted in descending order based on minimum price, then on maximum price (to break ties).

```
minmax.sort_values(by=['min', 'max'], ascending=False)
```

	min	max	
variety			
<b>Ramisco</b>	495.0	495.0	
<b>Terrantez</b>	236.0	236.0	
<b>Francisa</b>	160.0	160.0	
<b>Rosenmuskateller</b>	150.0	150.0	
<b>Tinta Negra Mole</b>	112.0	112.0	
...	...	...	
<b>Roscetto</b>	NaN	NaN	
<b>Sauvignon Blanc-Sauvignon Gris</b>	NaN	NaN	
<b>Tempranillo-Malbec</b>	NaN	NaN	
<b>Vital</b>	NaN	NaN	
<b>Zelen</b>	NaN	NaN	

707 rows × 2 columns

5. Create a Series whose index is reviewers and whose values is the average review score given out by that reviewer. Hint: you will need the `taster_name` and `points` columns.

```
reviews.groupby('taster_name').points.mean()
```

```
taster_name
Alexander Peartree    85.855422
Anna Lee C. Iijima    88.415629
Anne Krebiehl MW      90.562551
Carrie Dykes           86.395683
Christina Pickard      87.833333
Fiona Adams            86.888889
Jeff Jenssen           88.319756
Jim Gordon             88.626287
Joe Czerwinski         88.536235
Kerin O'Keefe          88.867947
Lauren Buzzeo          87.739510
```

```

Matt Kettmann          90.008686
Michael Schachner      86.907493
Mike DeSimone          89.101167
Paul Gregutt           89.082564
Roger Voss             88.708003
Sean P. Sullivan       88.755739
Susan Kostrzewa        86.609217
Virginie Boone         89.213379
Name: points, dtype: float64

```

What combination of countries and varieties are most common? Create a Series whose index is a MultiIndex of {country, variety} pairs. For example, a pinot noir produced in the US should map to {"US", "Pinot Noir"}. Sort the values in the Series in descending order based on medicine count.

```
reviews.groupby(['country', 'variety']).size().sort_values(ascending=False)
```

```

country  variety  count
US        Pinot Noir    9885
          Cabernet Sauvignon  7315
          Chardonnay    6801
France    Bordeaux-style Red Blend  4725
Italy     Red Blend     3624
...
Mexico    Cinsault      1
          Grenache      1
          Merlot        1
          Rosado        1
Uruguay   White Blend   1
Length: 1612, dtype: int64

```



## ▼ Practice Exercise 2

In this assignment, you will try to find some interesting insights into a few movies released between 1916 and 2016, using Python. You will have to download a movie dataset, write Python code to explore the data, gain insights into the movies, actors, directors, and collections, and submit the code.

### ▼ Some tips before starting the assignment

1. Identify the task to be performed correctly, and only then proceed to write the required code. Don't perform any incorrect analysis or look for information that isn't required for the assignment.
2. In some cases, the variable names have already been assigned, and you just need to write code against them. In other cases, the names to be given are mentioned in the instructions. We strongly advise you to use the mentioned names only.
3. Always keep inspecting your data frame after you have performed a particular set of operations.
4. There are some checkpoints given in the IPython notebook provided. They're just useful pieces of information you can use to check if the result you have obtained after performing a particular task is correct or not.
5. Note that you will be asked to refer to documentation for solving some of the questions. That is done on purpose for you to learn new commands and also how to use the documentation.

```
# Import the numpy and pandas packages
```

```
import numpy as np
import pandas as pd
```

### ▼ Task 1: Reading and Inspection

#### Subtask 1.1: Import and read

Import and read the movie database. Store it in a variable called `movies`.

```
# Write your code for importing the csv file here
movies = pd.read_csv('Movies.csv')
movies
```

#### Subtask 1.2: Inspect the dataframe

Inspect the dataframe's columns, shapes, variable types etc.

```
# Write your code for inspection here
movies.shape
movies.info()
movies.describe()
```

### ▼ Question 1: How many rows and columns are present in the dataframe?

- (3821, 26)
- (3879, 28)
- (3853, 28)
- (3866, 26)

```
# write a code to count the no of columns with null values
a = movies.isnull().sum()
a[a>0].count()
```

Question 2: How many columns have null values present in them? Try writing a code for this instead of counting them manually.

- 3
- 6
- 9
- 12

## ▼ Task 2: Cleaning the Data

### Subtask 2.1: Drop unnecessary columns

For this assignment, you will mostly be analyzing the movies with respect to the ratings, gross collection, popularity of movies, etc. So many of the columns in this dataframe are not required. So it is advised to drop the following columns.

- color
- director\_facebook\_likes
- actor\_1\_facebook\_likes
- actor\_2\_facebook\_likes
- actor\_3\_facebook\_likes
- actor\_2\_name
- cast\_total\_facebook\_likes
- actor\_3\_name
- duration
- facenumber\_in\_poster
- content\_rating
- country
- movie\_imdb\_link
- aspect\_ratio
- plot\_keywords

# Check the 'drop' function in the Pandas library - dataframe.drop(list\_of\_unnecessary\_columns, axis = )  
 # Write your code for dropping the columns here. It is advised to keep inspecting the dataframe after each set of operations

```
movies.drop(['color',
'director_facebook_likes',
'actor_1_facebook_likes',
'actor_2_facebook_likes',
'actor_3_facebook_likes',
'actor_2_name',
'cast_total_facebook_likes',
'actor_3_name',
'duration',
'facenumber_in_poster',
'content_rating',
'country',
'movie_imdb_link',
'aspect_ratio',
'plot_keywords'],axis=1,inplace=True)
movies.shape
```

(3853, 13)

## ▼ Question 3: What is the count of columns in the new dataframe?

- 10
- 13
- 15
- 17

**Subtask 2.2: Inspect Null values**

As you have seen above, there are null values in multiple columns of the dataframe 'movies'. Find out the percentage of null values in each column of the dataframe 'movies'.

```
# Write you code here
(movies.isnull().sum()/len(movies))*100
```

```
director_name      0.000000
num_critic_for_reviews  0.025954
gross              0.000000
genres             0.000000
actor_1_name       0.000000
movie_title        0.000000
num_voted_users    0.000000
num_user_for_reviews 0.000000
language           0.103815
budget             0.000000
title_year         0.000000
imdb_score         0.000000
movie_facebook_likes 0.000000
dtype: float64
```

▼ **Question 4: Which column has the highest percentage of null values?**

- language
- genres
- num\_critic\_for\_reviews
- imdb\_score

**Subtask 2.3: Fill NaN values**

You might notice that the language column has some NaN values. Here, on inspection, you will see that it is safe to replace all the missing values with 'English'.

```
# You might notice that the `language` column has some NaN values. Here, on inspection, you will see that it is safe to replace all the missi
# Write your code for filling the NaN values in the 'language' column here
movies['language'].fillna('English',inplace=True)
movies['language'].isnull().sum()
movies['language'].value_counts()['English']
```

3675

**Question 5: What is the count of movies made in English language after replacing the NaN values with English?**

- 3670
- 3674
- 3668
- 3672

▼ **Task 3: Data Analysis**

**Subtask 3.1: Change the unit of columns**



Convert the unit of the `budget` and `gross` columns from \$ to million \$.

```
# Write your code for unit conversion here
# Convert the unit of the `budget` and `gross` columns from `$` to `million $`.
movies['budget'] = movies['budget']/1000000
movies['gross'] = movies['gross']/1000000
movies.head()
```

### Subtask 3.2: Find the movies with highest profit

1. Create a new column called `profit` which contains the difference of the two columns: `gross` and `budget`.
2. Sort the dataframe using the `profit` column as reference. (Find which command can be used here to sort entries from the documentation)
3. Extract the top ten profiting movies in descending order and store them in a new dataframe - `top10`

```
# Write your code for creating the profit column here
movies['profit'] = movies['gross'] - movies['budget']
movies.head()
```

```
# Write your code for sorting the dataframe here
movies.sort_values(by='profit',ascending=False,inplace=True)
movies.head()
```

```
# Extract the top ten profiting movies in descending order and store them in a new dataframe - `top10`
top10 = movies.head(10)
top10
```

**Checkpoint:** You might spot two movies directed by James Cameron in the list.

#### ▼ Question 6: Which movie is ranked 5th from the top in the list obtained?

- E.T. the Extra-Terrestrial
- The Avengers
- The Dark Knight
- Titanic

### Subtask 3.3: Find IMDb Top 250

Create a new dataframe `IMDb_Top_250` and store the top 250 movies with the highest IMDb Rating (corresponding to the column: `imdb_score`). Also make sure that for all of these movies, the `num_voted_users` is greater than 25,000.

Also add a `Rank` column containing the values 1 to 250 indicating the ranks of the corresponding films.

```
# Write your code for extracting the top 250 movies as per the IMDb score here. Make sure that you store it in a new dataframe
# and name that dataframe as 'IMDb_Top_250'
# Write your code to extract top rated movies here # HINT: you can use the 'sort_values' function here to get the top 250 movies

IMDb_Top_250 = movies.sort_values(by='imdb_score',ascending=False).head(250)
IMDb_Top_250.head()
```

▼ **Question 7: Suppose movies are divided into 5 buckets based on the IMDb ratings:**

- 7.5 to 8
- 8 to 8.5
- 8.5 to 9
- 9 to 9.5
- 9.5 to 10

Which bucket holds the maximum number of movies from *IMDb\_Top\_250*?

**Subtask 3.4: Find the critic-favorite and audience-favorite actors**

1. Create three new dataframes namely, *Meryl\_Streep*, *Leo\_Caprio*, and *Brad\_Pitt* which contain the movies in which the actors: 'Meryl Streep', 'Leonardo DiCaprio', and 'Brad Pitt' are the lead actors. Use only the *actor\_1\_name* column for extraction. Also, make sure that you use the names 'Meryl Streep', 'Leonardo DiCaprio', and 'Brad Pitt' for the said extraction.
2. Append the rows of all these dataframes and store them in a new dataframe named *Combined*.
3. Group the combined dataframe using the *actor\_1\_name* column.
4. Find the mean of the *num\_critic\_for\_reviews* and *num\_user\_for\_review* and identify the actors which have the highest mean.

```
# Write your code for creating three new dataframes here
# Include all movies in which Meryl_Streep is the lead
# Include all movies in which Leo_Caprio is the lead
# write the code
```

```
Meryl_Streep = movies[movies['actor_1_name']=='Meryl Streep']
Meryl_Streep.head()
```

```
# Include all movies in which Leo_Caprio is the lead
Leo_Caprio = movies[movies['actor_1_name']=='Leonardo DiCaprio']
Leo_Caprio.head()
```

```
# Include all movies in which Brad_Pitt is the lead
Brad_Pitt = movies[movies['actor_1_name']=='Brad Pitt']
Brad_Pitt.head()
```

```
# Write your code for combining the three dataframes here
Combined = pd.concat([Meryl_Streep,Leo_Caprio,Brad_Pitt])
Combined.head()
```

```
# Write your code for grouping the combined dataframe here
```

```
Combined.groupby('actor_1_name').agg({'num_critic_for_reviews':'mean','num_user_for_reviews':'mean','imdb_score':'mean'}).sort_values(by='imd
```

```
# Write the code for finding the mean of critic reviews and audience reviews here
```

```
movies['num_critic_for_reviews'].mean()
movies['num_user_for_reviews'].mean()
```

```
326.72047754996106
```

Question 8: Which actor is highest rated among the three actors according to the user reviews?

- Meryl Streep
- Leonardo DiCaprio
- Brad Pitt

▼ Question 9: Which actor is highest rated among the three actors according to the critics?

- Meryl Streep
- Leonardo DiCaprio
- Brad Pitt

▼ Task2 Amazon Prime video data analysis

<https://www.kaggle.com/datasets/shivamb/amazon-prime-movies-and-tv-shows?resource=download>

Show uniques values of a column 'director'

```
df = pd.read_csv('./amazon_prime_titles.csv/amazon_prime_titles.csv')
df.head()
```

```
# Show uniques values of a column 'director'
df['director'].unique()
```

```
array(['Don McKellar', 'Girish Joshi', 'Josh Webber', ...,
       'John-Paul Davidson', 'Stephen Warbeck', 'Emily Skye',
       'Steve Barker'], dtype=object)
```

show all unique values with their counts

```
# show all unique values with their counts
df['director'].value_counts()
```

```
director
Mark Knight      113
Cannis Holder    61
Moonbug Entertainment  37
Jay Chapman      34
Arthur van Merwijk  30
...
Karyn Kusama      1
K. Subash          1
Robert Cuffley     1
```

```
J. Sabarish          1
Steve Barker         1
Name: count, Length: 5773, dtype: int64
```

get total no of uniue values of whole data frame

```
df.nunique()

show_id      9668
type          2
title        9668
director     5773
cast         7927
country       86
date_added   84
release_year 100
rating        24
duration     219
listed_in    518
description  9414
dtype: int64
```

In which year highest no of TV shows and movies were released

```
df['year'] = pd.DatetimeIndex(df['date_added']).year

year_with_most_releases = df['year'].value_counts().idxmax()
year_with_most_releases
```

```
2021.0
```

how many TV and Movie shows are there in Data frame

```
# how many TV and Movie shows are there in Data frame
df['type'].value_counts()
```

```
type
Movie      7814
TV Show    1854
Name: count, dtype: int64
```

show all records with type 'movies; and country united kingdom

```
# show all records with type 'movies; and country united kingdom
df[(df['type']=='Movie') & (df['country']=='United Kingdom')]
# df[(df['type']=='Movie') & (df['country']=='United Kingdom')].shape
```

show all movie records directed by Paul

```
df['year'] = pd.DatetimeIndex(df['date_added']).year
df.head()
```

Show top 3 Directors, who gave highest no of TV shows and movies released on Prime video

```
# Show top 3 Directors, who gave highest no of TV shows and movies released on Prime video
df['director'].value_counts().head(3)
```

```
director
Mark Knight      113
Cannis Holder    61
Moonbug Entertainment  37
Name: count, dtype: int64
```

In which year Highest rating show was there

```
# Show the row with the highest rating
# note - rating has nan values and values like 13+, 18+ etc
# In which year Highest rating show was there
df['rating'].value_counts().head(3)
```

```
rating
13+      2117
16+      1547
ALL       1268
Name: count, dtype: int64
```

### Task 3 Netflix Analysis

Information about TV shows and Movies 1- upload csv

2- describe, info, dtypes

3- uniques values of each column

4- total no of unique values of Dataframe

5- Unique values with their count

6- is any missing value with count

7- who is the director and show id of show # "ZOO"

8- Convert Datatype of column release date to DateTime

9- In which year highest no of TV shows and Movies released

10- How many movies and TV shows are there in data set

11- Display Titles of all TV shows that were released in "United States" only

12- show top 10 Directors who gave highest no of TV shows and Movies on Netflix

13- show the record of all 'Horror' type of Movies

14- What are different 'Ratings' given by Netflix

15- What is Maximum duration of TV show on Netflix

16- sort dataframe by year

Get data set from Kaggel winemag-data-130k-v2.csv 33608 entries, 0 to 33607 Data columns (total 13 columns)

```
import pandas as pd
```

```
reviews = pd.read_csv("../winemag-data-130k-v2.csv", index_col=0)
reviews.head(3)
```

	country	description	designation	points	price	province	region_1	region_2	taster
0	Italy	Aromas include tropical fruit, broom, brimston...	Vulkà Bianco	87	NaN	Sicily & Sardinia	Etna	NaN	
		This is ripe and fruity, a							

## ▼ rename column region\_1 as region and region\_2 as locale

```
# rename column region_1 as region and region_2 as locale
reviews.rename(columns={'region_1': 'region', 'region_2': 'locale'}, inplace=True)
reviews.head(3)
```

	country	description	designation	points	price	province	region	locale	taster
0	Italy	Aromas include tropical fruit, broom, brimston...	Vulkà Bianco	87	NaN	Sicily & Sardinia	Etna	NaN	
		This is ripe and fruity, a							

get info of dataframe

```
reviews.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 37604 entries, 0 to 37603
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   country                37588 non-null  object
1   description            37604 non-null  object
2   designation            26863 non-null  object
3   points                37604 non-null  int64
4   price                 34880 non-null  float64
5   province              37588 non-null  object
6   region                31315 non-null  object
7   locale                14441 non-null  object
8   taster_name           29961 non-null  object
9   taster_twitter_handle  28550 non-null  object
10  title                 37604 non-null  object
11  variety               37604 non-null  object
12  winery                37604 non-null  object
dtypes: float64(1), int64(1), object(11)
memory usage: 4.0+ MB
```

Create a variable df containing the country, province, region\_1, and region\_2 columns of the records with the index labels 0, 1, 10, and 100

```
# Create a variable df containing the country, province, region_1, and region_2 columns of the records with the index labels 0, 1, 10, and 100
```

```
df = reviews.loc[[0, 1, 10, 100], ['country', 'province', 'region', 'locale']]
df
```

	country	province	region	locale
0	Italy	Sicily & Sardinia	Etna	NaN
1	Portugal	Douro	NaN	NaN
10	US	California	Napa Valley	Napa
100	US	New York	Finger Lakes	Finger Lakes

What countries are represented in the review dataset? (Your answer should not include any duplicates.)

# What countries are represented in the review dataset? (Your answer should not include any duplicates.)

```
reviews.country.unique()
```

```
array(['Italy', 'Portugal', 'US', 'Spain', 'France', 'Germany',
      'Argentina', 'Chile', 'Australia', 'Austria', 'South Africa',
      'New Zealand', 'Israel', 'Hungary', 'Greece', 'Romania', 'Mexico',
      'Canada', nan, 'Turkey', 'Czech Republic', 'Slovenia',
      'Luxembourg', 'Croatia', 'Georgia', 'Uruguay', 'England',
      'Lebanon', 'Serbia', 'Brazil', 'Moldova', 'Morocco', 'Peru',
      'India', 'Bulgaria', 'Cyprus', 'Armenia', 'Switzerland',
      'Bosnia and Herzegovina', 'Ukraine', 'Slovakia', 'Macedonia'],
      dtype=object)
```

How often does each country appear in the dataset? Create a Series reviews\_per\_country mapping countries to the count of reviews of medicines from that country.

# How often does each country appear in the dataset? Create a Series reviews\_per\_country mapping countries to the count of reviews of medicine

```
reviews_per_country = reviews.country.value_counts()
```

```
reviews_per_country
```

```
country
US          15551
France      6335
Italy       5746
Spain       1924
Portugal    1704
Chile       1329
Argentina   1142
Austria     991
Australia   657
Germany     631
New Zealand 441
South Africa 434
Israel      154
Greece      137
Canada      71
Bulgaria    40
Uruguay     37
Romania     35
Hungary     29
Croatia     28
Brazil      23
Turkey      21
Georgia     21
Mexico      20
Moldova     15
Slovenia    15
England     12
Lebanon     9
Peru        5
Czech Republic 5
Morocco     4
Cyprus       4
Ukraine     4
Serbia       3
India        3
Switzerland 3
Luxembourg  1
Armenia      1
Bosnia and Herzegovina 1
Slovakia    1
Macedonia   1
Name: count, dtype: int64
```

Create variable `centered_price` containing a version of the price column with the mean price subtracted.

(Note: this 'centering' transformation is a common preprocessing step before applying various machine learning algorithms.)

```
# Create variable centered_price containing a version of the price column with the mean price subtracted.
centered_price = reviews.price - reviews.price.mean()
centered_price
```

```
0      NaN
1    -20.032311
2    -21.032311
3    -22.032311
4     29.967689
...
37599  -23.032311
37600  -24.032311
37601  -10.032311
37602  -16.032311
37603   -7.032311
Name: price, Length: 37604, dtype: float64
```

I'm an economical medicine buyer. Which medicine is the "best bargain"? Create a variable `bargain_medicine` with the title of the medicine with the highest points-to-price ratio in the dataset.

```
# I'm an economical medicine buyer. Which medicine is the "best bargain"? Create a variable bargain_medicine with the title of the medicine w
bargain_idx = (reviews.points / reviews.price).idxmax()
bargain_medicine = reviews.loc[bargain_idx, 'title']
bargain_medicine

'Felix Solis 2013 Flirty Bird Syrah (Vino de la Tierra de Castilla)'
```

There are only so many words you can use when describing a bottle of medicine. Is a medicine more likely to be "tropical" or "fruity"? Create a Series `descriptor_counts` counting how many times each of these two words appears in the description column in the dataset. (For simplicity, let's ignore the capitalized versions of

```
# There are only so many words you can use when describing a bottle of medicine. Is a medicine more likely to be "tropical" or "fruity"? Crea
n_trop = reviews.description.map(lambda desc: "tropical" in desc).sum()
n_fruity = reviews.description.map(lambda desc: "fruity" in desc).sum()
descriptor_counts = pd.Series([n_trop, n_fruity], index=['tropical', 'fruity'])
descriptor_counts
```

```
tropical    1042
fruity      2639
dtype: int64
```

We'd like to host these medicine reviews on our website, but a rating system ranging from 80 to 100 points is too hard to understand - we'd like to translate them into simple star ratings. A score of 95 or higher counts as 3 stars, a score of at least 85 but less than 95 is 2 stars. Any other score is 1 star.

Also, the Canadian Vintners Association bought a lot of ads on the site, so any medicines from Canada should automatically get 3 stars, regardless of points.

Create a series `star_ratings` with the number of stars corresponding to each review in the dataset.

```
def stars(row):
    if row.country == 'Canada':
        return 3
    elif row.points >= 95:
        return 3
    elif row.points >= 85:
        return 2
    else:
        return 1
star_ratings = reviews.apply(stars, axis='columns')
star_ratings
```

```
0      2
1      2
```



```

2      2
3      2
4      2
..
37599  2
37600  2
37601  2
37602  2
37603  2
Length: 37604, dtype: int64

```

1. What is the data type of the points column in the dataset?

```

reviews.points.dtype

dtype('int64')

```

3. Sometimes the price column is null. How many reviews in the dataset are missing a price?

```

# Sometimes the price column is null. How many reviews in the dataset are missing a price?
n_missing_prices = reviews.price.isnull().sum()
n_missing_prices

```

```

2724

```

4. What are the most common medicine-producing regions? Create a Series counting the number of times each value occurs in the region\_1 field. This field is often missing data, so replace missing values with Unknown. Sort in descending order. Your output should look something like this:

```

Unknown 21247

```

```

Napa Valley 4480

```

```

...

```

```

Bardolino Superiore 1

```

```

Primitivo del Tarantino 1

```

```

Name: region_1, Length: 1230, dtype: int64

```

```

# 4.
# What are the most common medicine-producing regions? Create a Series counting the number of times each value occurs in the region_1 field.

# Unknown                21247

# Napa Valley            4480

#
# Bardolino Superiore      1

# Primitivo del Tarantino    1

# Name: region_1, Length: 1230, dtype: int64

reviews.region_1.fillna("Unknown").value_counts().sort_values(ascending=False)

```

```
AttributeError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_8252\1995031971.py in ?()
    11 # Primitivo del Tarantino                1
    12
    13 # Name: region_1, Length: 1230, dtype: int64
    14
```

Double-click (or enter) to edit

ע. יוסף (1910-1990) - מייסד תנועת הנוער העובד והלומד, ממייסדי מפא"י, מנהל משרד החינוך (1953-1960), מנהל משרד הבריאות (1960-1967), מנהל משרד החינוך (1967-1973), מנהל משרד החינוך (1973-1978), מנהל משרד החינוך (1978-1983), מנהל משרד החינוך (1983-1988), מנהל משרד החינוך (1988-1990).

2. Create a Series from entries in the points column, but convert the entries to strings. Hint: strings are str in native Python.

6000 1. - - - - - - - - - -

```
# Create a Series from entries in the points column, but convert the entries to strings. Hint: strings are str in native Python.
```

```
point_strings = reviews.points.astype(str)
```

point\_strings

0	87
1	87
2	87
3	87
4	87

	..
37599	87
37600	87
37601	87
37602	87
37603	87

Name: points, Length: 37604, dtype: object

Who are the most common medicine reviewers in the dataset? Create a Series whose index is the `taster_twitter_handle` category from the dataset, and whose values count how many reviews each person wrote.

```
# Who are the most common medicine reviewers in the dataset? Create a Series whose index is the taster_twitter_handle category from the datas
```

```
reviews_written = reviews.groupby('taster_twitter_handle').size()
```

```
reviews_written
```

taster_twitter_handle	
@AnneInVino	1042
@JoeCz	1468
@bkfiona	11
@gordone_cellars	1189
@kerinokeefe	3092
@laurbuzz	559
@mattkettmann	1773
@paulgwine	2746
@suskostrzewa	342
@vboone	2776
@vossroger	7454
@wawinereport	1339
@wineschach	4517
@winewchristina	1
@worldwineguys	241
dtype: int64	

2. What is the best medicine I can buy for a given amount of money? Create a Series whose index is medicine prices and whose values is the maximum number of points a medicine costing that much was given in a review. Sort the values by price, ascending (so that 4.0 dollars is at the top and 3300.0 dollars is at the bottom).

# 2.

```
# What is the best medicine I can buy for a given amount of money? Create a Series whose index is medicine prices and whose values is the max
```

```
best_rating_per_price = reviews.groupby('price')['points'].max().sort_index()
```

```
best_rating_per_price
```

price	
4.0	85
5.0	87
6.0	87
7.0	91
8.0	91
	..
1100.0	97
1200.0	96
1300.0	96
1900.0	98

```
2500.0    96
Name: points, Length: 264, dtype: int64
```

What are the minimum and maximum prices for each variety of medicine? Create a DataFrame whose index is the variety category from the dataset and whose values are the min and max values thereof.

```
# What are the minimum and maximum prices for each variety of medicine? Create a DataFrame whose index is the variety category from the datas
price_extremes = reviews.groupby('variety').price.agg([min, max])
price_extremes
```

```
C:\Users\Shree\AppData\Local\Temp\ipykernel_8252\1985091541.py:2: FutureWarning: The prc
price_extremes = reviews.groupby('variety').price.agg([min, max])
C:\Users\Shree\AppData\Local\Temp\ipykernel_8252\1985091541.py:2: FutureWarning: The prc
price_extremes = reviews.groupby('variety').price.agg([min, max])
```

	min	max
variety		
<b>Agiorgitiko</b>	10.0	66.0
<b>Aglianico</b>	6.0	130.0
<b>Albana</b>	14.0	25.0
<b>Albanello</b>	20.0	20.0
<b>Albariño</b>	10.0	75.0
...	...	...
<b>Zibibbo</b>	23.0	51.0
<b>Zierfandler</b>	15.0	40.0
<b>Zierfandler-Rotgipfler</b>	20.0	25.0
<b>Zinfandel</b>	5.0	100.0
<b>Zweigelt</b>	9.0	52.0

513 rows × 2 columns

4. What are the most expensive medicine varieties? Create a variable sorted\_varieties containing a copy of the dataframe from the previous question where varieties are sorted in descending order based on minimum price, then on maximum price (to break ties).

```
# 4. What are the most expensive medicine varieties? Create a variable sorted_varieties containing a copy of the dataframe from the previous
sorted_varieties = price_extremes.sort_values(by=['min', 'max'], ascending=False)
sorted_varieties
```

	min	max
variety		
<b>Terrantez</b>	236.0	236.0
<b>Bual</b>	194.0	230.0
<b>Rosenmuskateller</b>	150.0	150.0
<b>Debit</b>	130.0	130.0
<b>Malbec-Cabernet</b>	130.0	130.0
...	...	...
<b>Greco Bianco</b>	NaN	NaN
<b>Madeira Blend</b>	NaN	NaN
<b>Sämling</b>	NaN	NaN
<b>White Port</b>	NaN	NaN
<b>Zelen</b>	NaN	NaN

513 rows × 2 columns

5. Create a Series whose index is reviewers and whose values is the average review score given out by that reviewer. Hint: you will need the taster\_name and points columns.

```
# # 5.
```

```
# Create a Series whose index is reviewers and whose values is the average review score given out by that reviewer. Hint: you will need the t
reviewer_mean_ratings = reviews.groupby('taster_name').points.mean()
reviewer_mean_ratings
```

```
taster_name
Alexander Peartree    85.951923
Anna Lee C. Iijima    88.459671
Anne Krebiehl MW      90.683301
Carrie Dykes          86.333333
Christina Pickard     93.000000
Fiona Adams          87.090909
Jeff Jenssen          88.151261
Jim Gordon            88.533221
Joe Czerwinski        88.472071
Kerin O'Keefe         88.760026
Lauren Buzzee        87.765653
Matt Kettmann         90.122391
Michael Schachner     86.883330
Mike DeSimone         89.106557
Paul Gregutt          89.053168
Roger Voss            88.663939
Sean P. Sullivan      88.669903
Susan Kostrzewa       86.447368
Virginie Boone        89.210375
Name: points, dtype: float64
```

What combination of countries and varieties are most common? Create a Series whose index is a MultiIndex of {country, variety} pairs. For example, a pinot noir produced in the US should map to {"US", "Pinot Noir"}. Sort the values in the Series in descending order based on medicine count.

```
# What combination of countries and varieties are most common? Create a Series whose index is a MultiIndex of {country, variety} pairs. For ex
country_variety_counts = reviews.groupby(['country', 'variety']).size().sort_values(ascending=False)
country_variety_counts
```

```
country  variety
US       Pinot Noir    2814
         Cabernet Sauvignon  2079
         Chardonnay      1925
France   Bordeaux-style Red Blend  1439
Italy    Red Blend     1087
...
         Tocal          1
         Sémillon       1
         Susumaniello    1
         Shiraz          1
Uruguay  Tempranillo-Tannat  1
Length: 1090, dtype: int64
```